

T04

# CPU Virtualization An Introduction to Scheduling

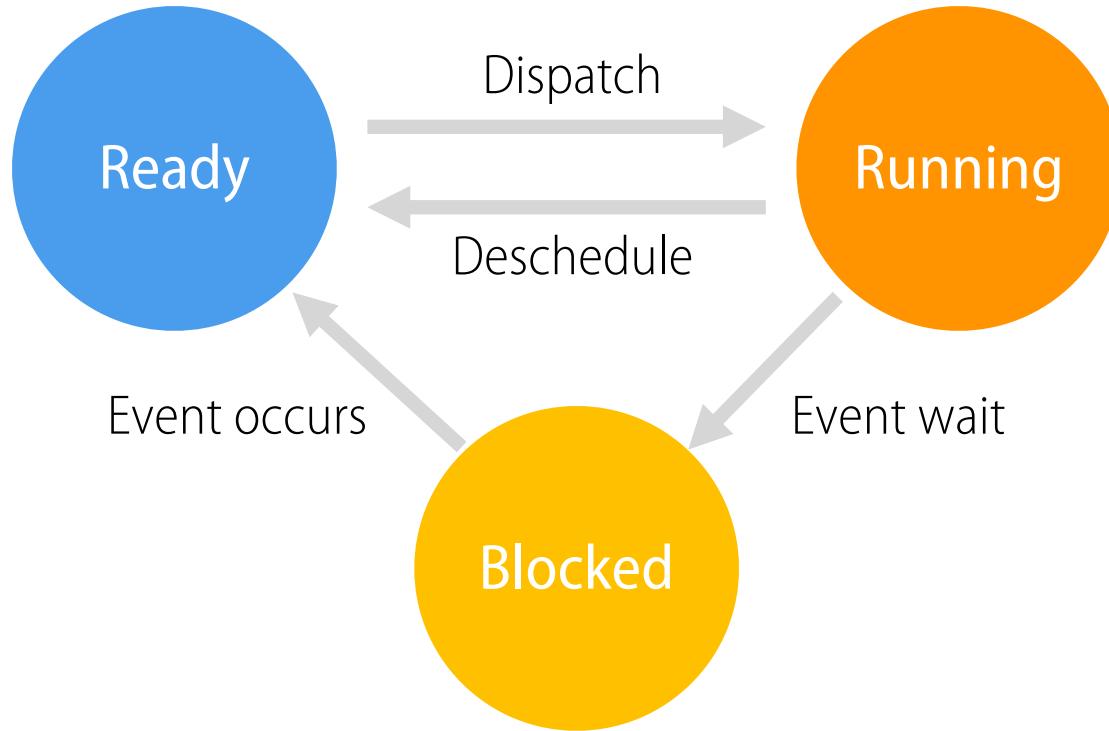
***Disclaimer:** This lecture slide set was developed to support the MC504 Operating System course at the University of Campinas and is mainly based on Chapter 07 of the Operating Systems: Three Easy Pieces book by Remzi and Andrea Arpaci-Dusseau of the University of Wisconsin, available at [pages.cs.wisc.edu/~remzi/OSTEP/](http://pages.cs.wisc.edu/~remzi/OSTEP/).*

# Two Components

- Dispatcher (Previous lecture)
  - Low-level mechanism
  - Performs context-switch
    - Switch from user mode to kernel mode
    - Save execution state (registers) of old process in PCB
    - Insert PCB in ready queue
    - Load state of next process from PCB to registers
    - Switch from kernel to user mode
    - Jump to instruction in new user process
- Scheduler (Today)
  - Applies a policy to determine which process gets CPU when there is a process switch.

Review

# State transitions



How to transition? ("mechanism")  
When to transition? ("policy")

# Usual Scheduling Performance Metrics

- Minimize turnaround time
  - Do not want to wait long for job to complete
  - $T_{turnaround} = T_{completion} - T_{arrival}$
- Minimize response time
  - Schedule interactive jobs promptly so users see output quickly
  - $T_{response} = T_{initial\_schedule} - T_{arrival}$
- Minimize waiting time
  - Do not spend much time in Ready queue
- Maximize throughput
  - Want many jobs to complete per unit of time
- Maximize resource utilization
  - Keep expensive devices busy
- Minimize overhead
  - Reduce number of context switches
- Maximize fairness
  - Give all jobs same amount of CPU over some time interval

# Scheduling basics

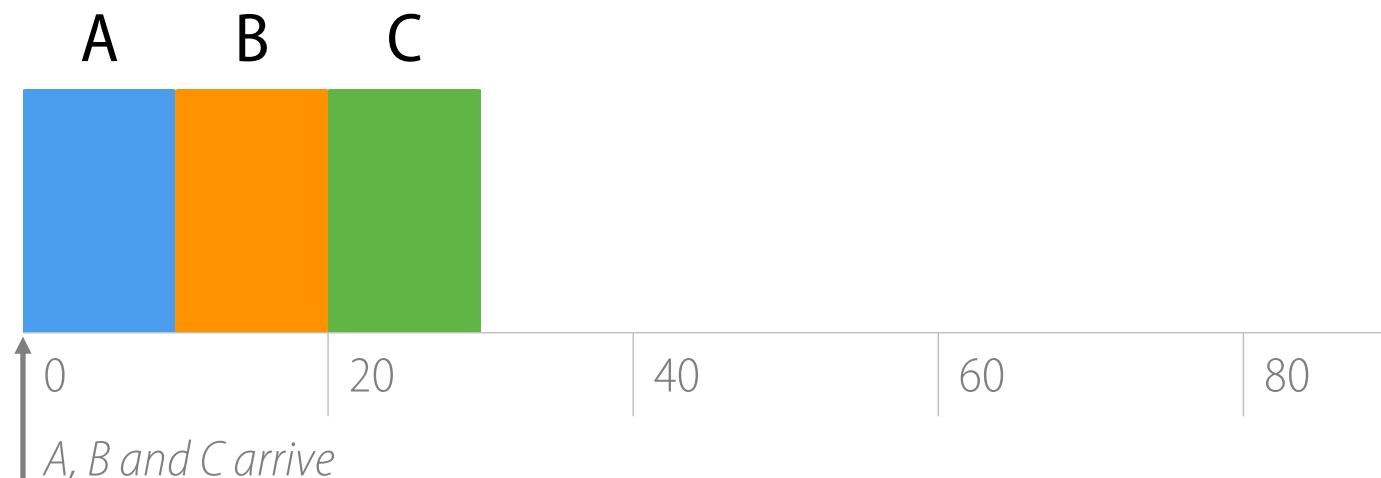
Workload	Scheduler	Metrics
$T_{arrival}$	FIFO	$T_{turnaround}$
$T_{run}$	SJF	$T_{response}$
	STCF	
	RR	

# Workload Assumptions

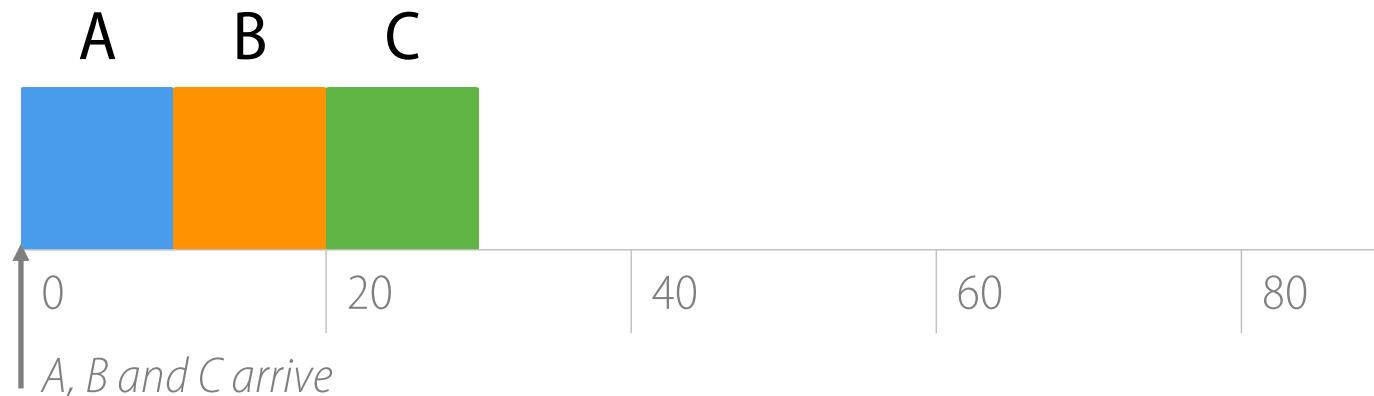
1. All jobs run for the same amount of time.
2. All jobs arrive at the same time.
3. All jobs only use the CPU (i.e. no I/O).
4. Run-time of each job is known.

# First In, First Out (FIFO)

- Also called First Come, First Served (FCFS)
  - Very simple and easy to implement
  - Jobs are executed in  $T_{arrival}$  order
- Example:
  - Job A arrives just before job B which arrives just before job C.
  - Each job runs for 10 seconds.



# FIFO (identical jobs)



- What is the average turnaround time?
  - $T_{turnaround} = T_{completion} - T_{arrival}$
  - $aT_{turnaround} = 10 - 0$
  - $bT_{turnaround} = 20 - 0$
  - $cT_{turnaround} = 30 - 0$
  - $avgT_{turnaround} = \frac{10+20+30}{3} = 20$

Why wouldn't FIFO be  
our solution of choice?

# Relaxed Workload Assumptions

- ~~1. All jobs run for the same amount of time.~~
2. All jobs arrive at the same time.
3. All jobs only use the CPU (i.e. no I/O).
4. Run-time of each job is known.

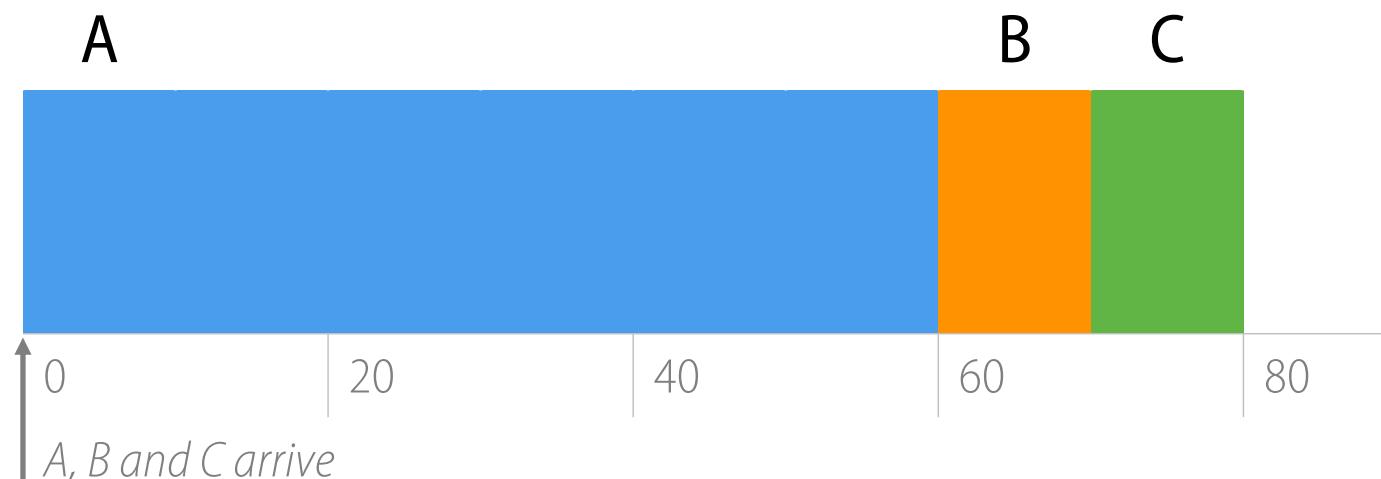
# Any problematic workloads for FIFO?

Workload	Scheduler	Metrics
$T_{arrival}$	FIFO	$T_{turnaround}$
$T_{run}$	SJF	$T_{response}$
	STCF	
	RR	

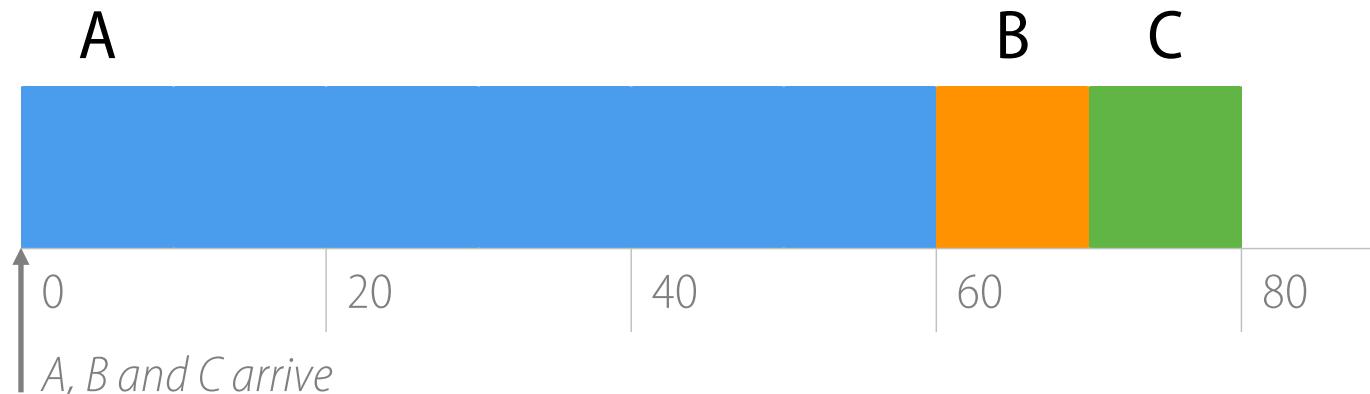
- Think of a really big first job...
  - What would happen to FIFO's turnaround time?

# FIFO (big first job)

- Workload
  - Job A arrives just before job B which arrives just before job C.
  - Job A runs for 60 seconds, jobs B and C run for 10 seconds.



# FIFO (identical jobs)



- What is the average turnaround time?
  - $T_{turnaround} = T_{completion} - T_{arrival}$
  - $aT_{turnaround} = 60 - 0$
  - $bT_{turnaround} = 70 - 0$
  - $cT_{turnaround} = 80 - 0$
  - $avgT_{turnaround} = \frac{60+70+80}{3} = 70$

# The Convoy Effect



© Getty

Jake Fear, of Highbridge, Somerset, was fined £190 after driving at 25mph for more than three miles on the A39 near Glastonbury in Somerset (file photo)

10 +3

# Overtaking the tractor

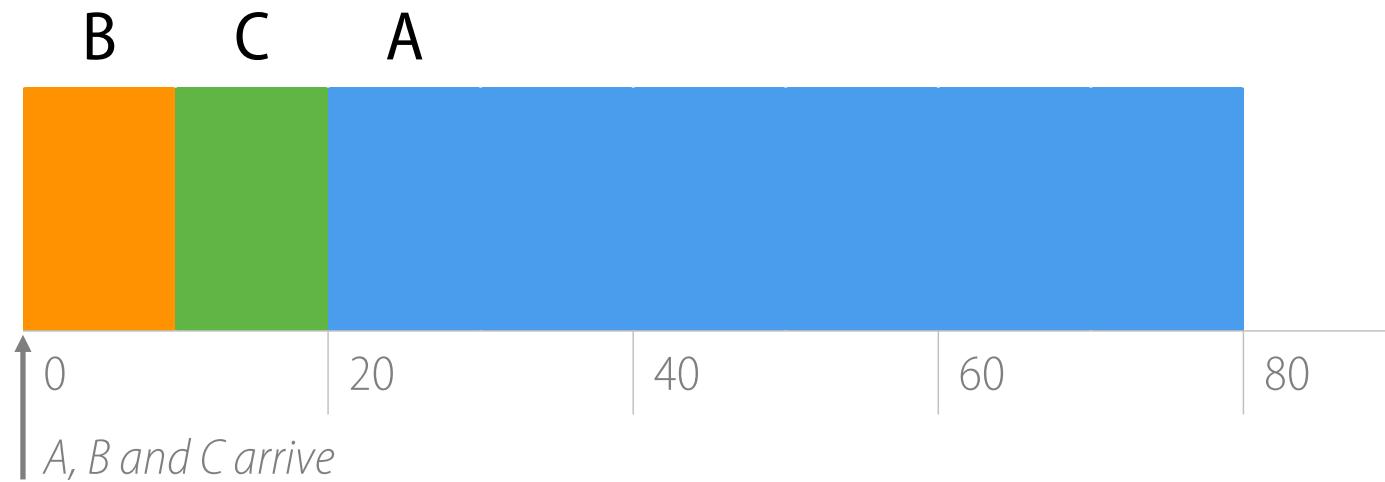
- What was the problem with FIFO?
  - Turnaround time can suffer when short jobs must wait for long jobs
- Let's try a new scheduler: **SJF (Shortest Job First)**
  - Choose job with smallest run time

# Scheduling basics

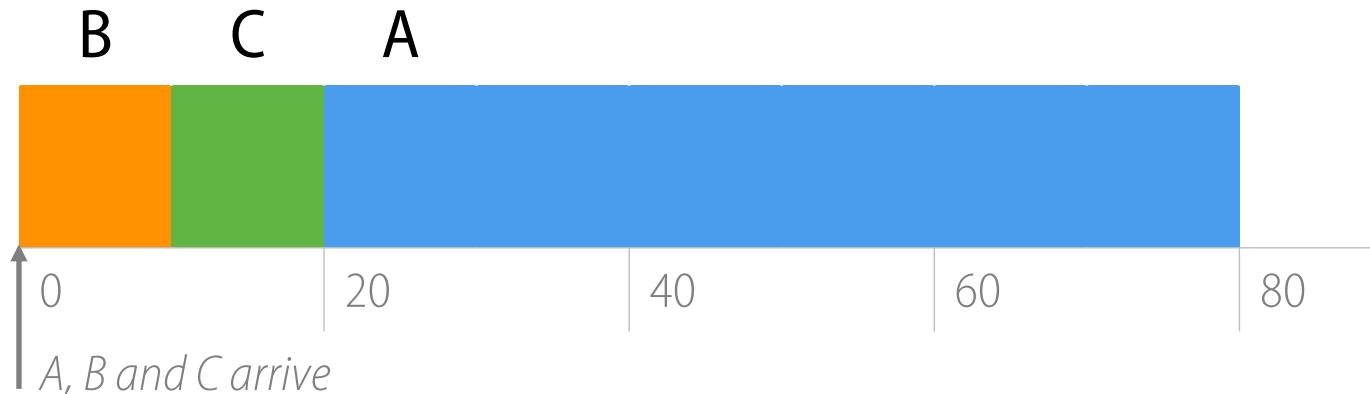
Workload	Scheduler	Metrics
$T_{arrival}$	FIFO	$T_{turnaround}$
$T_{run}$	SJF	$T_{response}$
	STCF	
	RR	

# SJF (shortest job first)

- Workload
  - Jobs A, B and C arrive at the same time.
  - Job A runs for 60 seconds, jobs B and C run for 10 seconds.



# SJF (shortest job first)



- What is the average turnaround time?
  - $T_{turnaround} = T_{completion} - T_{arrival}$
  - $aT_{turnaround} = 80 - 0$
  - $bT_{turnaround} = 10 - 0$
  - $cT_{turnaround} = 20 - 0$
  - $avgT_{turnaround} = \frac{80+10+20}{3} = 37$

# SJF (shortest job first)

- SJF is provably optimal, for minimizing average turnaround time (with no preemption).
  - Scheduling a shorter job before a longer job improves the turnaround time of the shorter job more than it harms the turnaround time of the longer job.

Why wouldn't SJF be  
our solution of choice?

# Relaxed Workload Assumptions

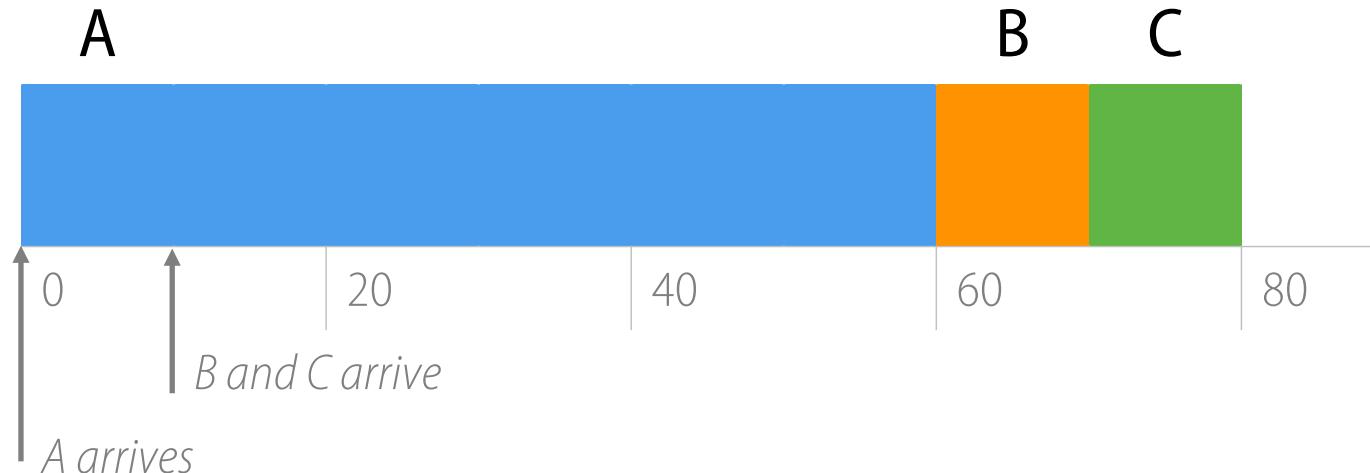
- ~~1. All jobs run for the same amount of time.~~
- ~~2. All jobs arrive at the same time.~~
3. All jobs only use the CPU (i.e. no I/O).
4. Run-time of each job is known.

# Any problematic workloads for SJF?

Workload	Scheduler	Metrics
$T_{arrival}$	FIFO	$T_{turnaround}$
$T_{run}$	SJF	$T_{response}$
	STCF	
	RR	

- Imagine that the shorter jobs arrive after a longer one...
  - Let's make  $aT_{arrival} = 0$  and  $bT_{arrival} = cT_{arrival} = 10$ .
  - What would happen to SJF's turnaround time?

# Stuck behind a tractor again...

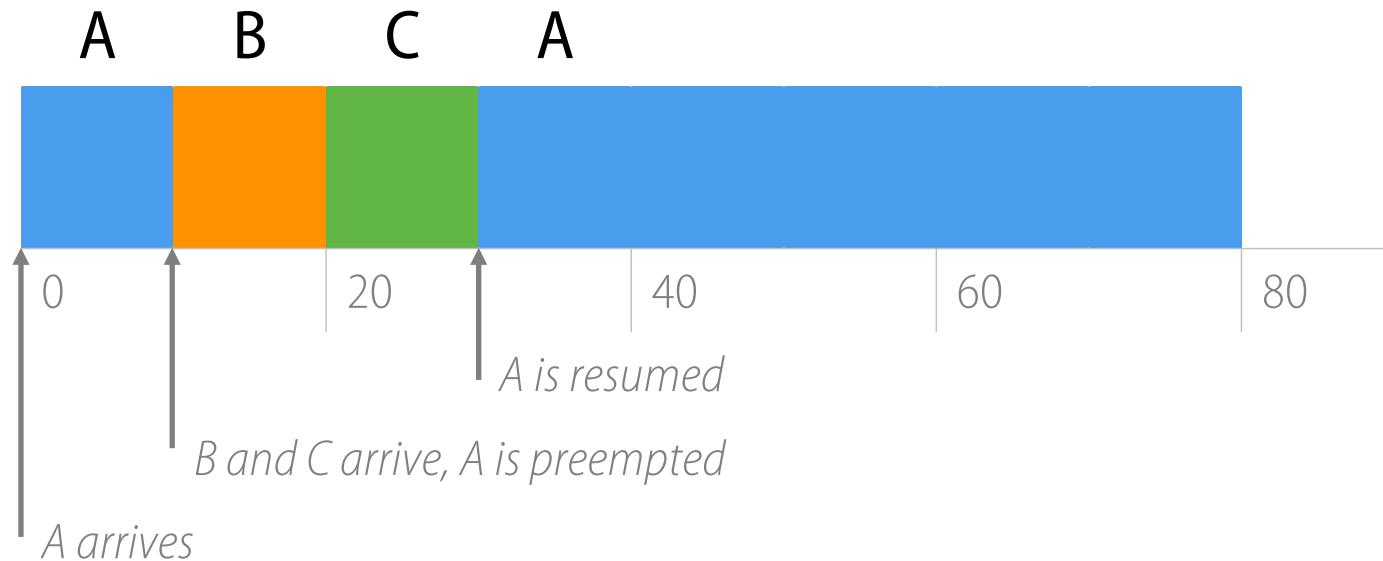


- What is the average turnaround time?
  - $T_{turnaround} = T_{completion} - T_{arrival}$
  - $aT_{turnaround} = 60 - 0$
  - $bT_{turnaround} = 70 - 10$
  - $cT_{turnaround} = 80 - 10$
  - $avgT_{turnaround} = \frac{60+60+70}{3} = 63$

# A case for preemptive scheduling

- What was the problem with FIFO and SJF?
  - They are non-preemptive, i.e. they only schedule a new job when the previous job voluntarily relinquishes the CPU (performs I/O or exits)
- Let's try a new scheduler: **STCF (Shortest Time-to-Completion First)**
  - Also known as Preemptive Shortest Job First (PSJF) and Shortest Remaining-Time First (SRTF).
  - Potentially schedule a different job at any point by taking CPU away from running job.
  - Always run the quickest-to-complete job.

# Preemptive STCF: overtaking the tractor again



- What is the average turnaround time?
  - $T_{turnaround} = T_{completion} - T_{arrival}$
  - $aT_{turnaround} = 80 - 0$
  - $bT_{turnaround} = 20 - 10$
  - $cT_{turnaround} = 30 - 10$
  - $avgT_{turnaround} = \frac{80+10+20}{3} = 37$

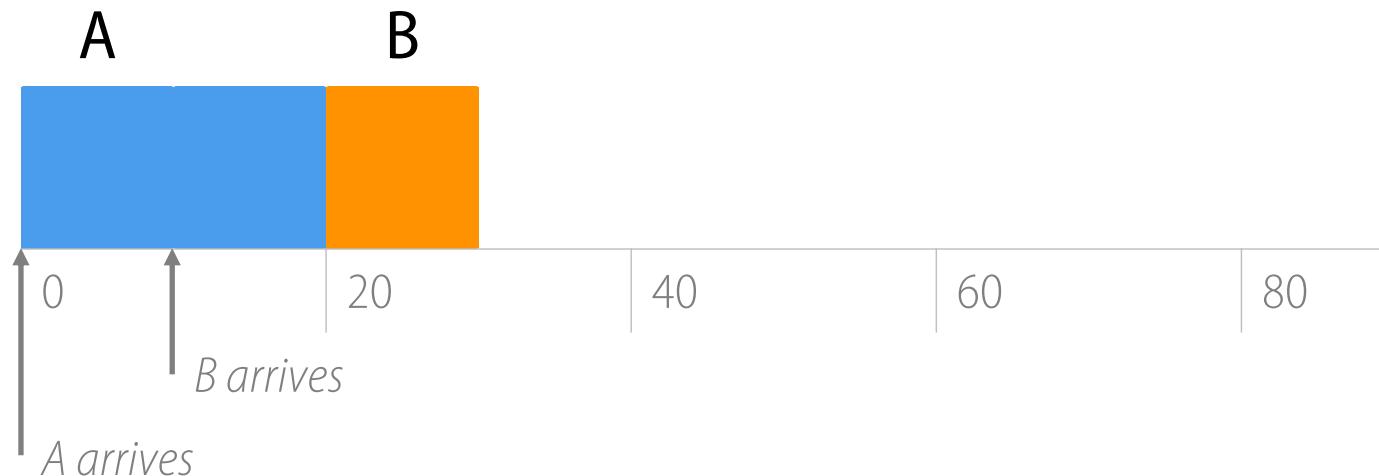
Why wouldn't STCF be  
our solution of choice?

# Response time

- Sometimes care about when a job starts instead of when it finishes...
- Response time
  - Time until the job is scheduled for the first time
  - $T_{response} = T_{firstrun} - T_{arrival}$
- STCF and related techniques are not particularly good for response time.
- **How can we build a scheduler that is sensitive to response time?**

# Response vs. Turnaround

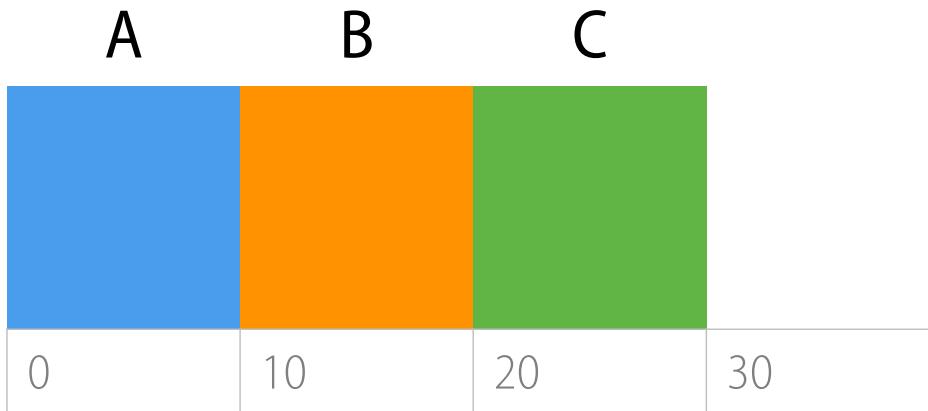
- In the non-preemptive example below...
  - B's turnaround time is...
    - $T_{turnaround} = T_{completion} - T_{arrival} = 30 - 10 = 20$
  - B's response time is...
    - $T_{response} = T_{firstrun} - T_{arrival} = 20 - 10 = 10$



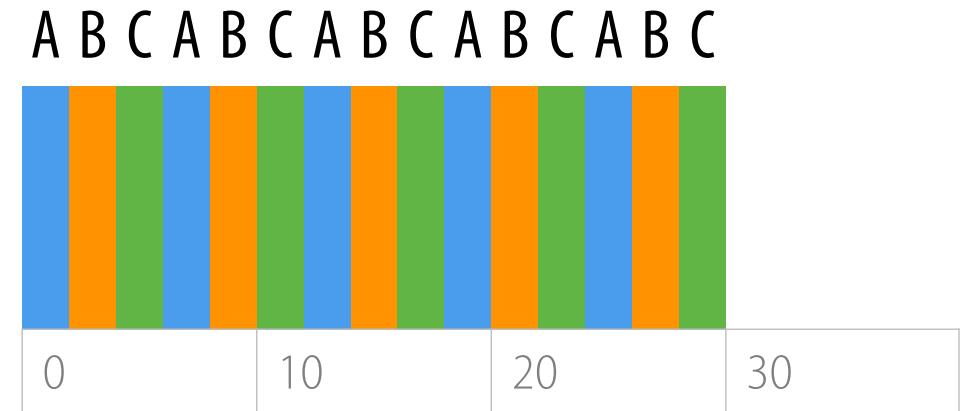
# Caring about response time

- What was the problem with FIFO, SJF and STCF?
  - They may show poor response time.
- Let's try a new scheduler: **RR (Round Robin)**
  - Run a job for a **time slice** and then switch to the next job in the ready queue until all jobs are finished.
    - Time slice is sometimes called a *scheduling quantum*.
  - The length of the time slice must be *a multiple of* the timer-interrupt period.

# FIFO vs. RR (simultaneous identical jobs)



$$T_{response} = (0 + 10 + 20)/3 = 10$$

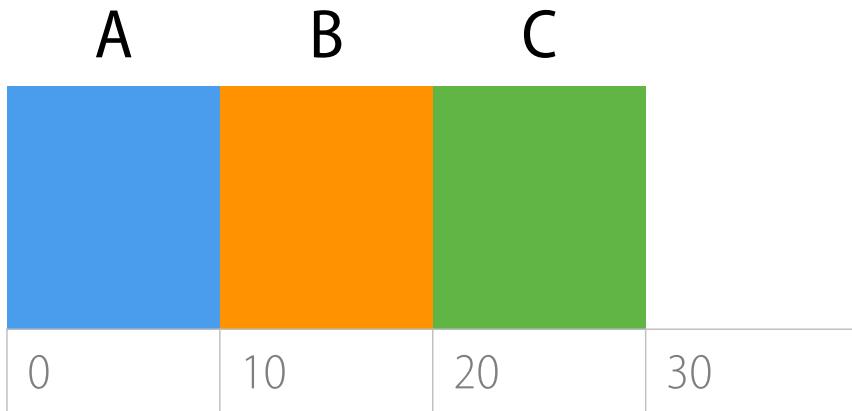


$$T_{response} = (0 + 2 + 4)/3 = 2$$

Why wouldn't RR be  
our solution of choice?

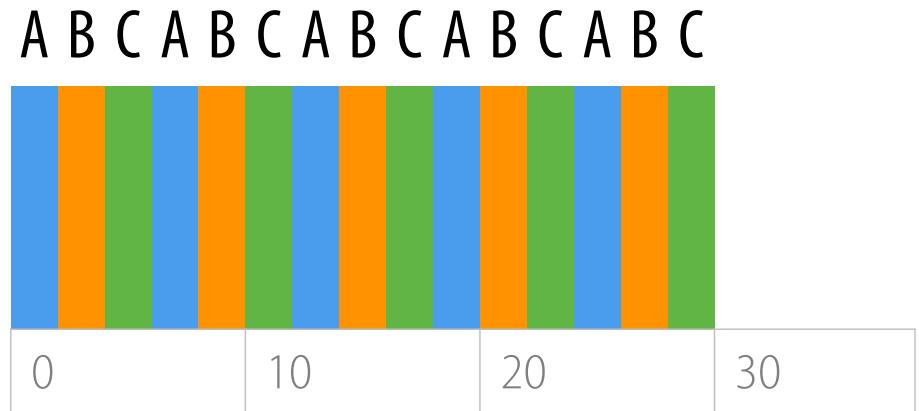
# In what ways isn't RR ideal?

- RR is fair, but performs poorly on metrics such as turnaround time.
  - Let's calculate the average turnaround time for both options in our last example.



$$avgT_{response} = (0 + 10 + 20)/3 = 10$$

$$\begin{aligned} avgT_{turnaround} &= \\ &= ((10 - 0) + (20 - 0) + (30 - 0))/3 = 20 \end{aligned}$$



$$avgT_{response} = (0 + 2 + 4)/3 = 2$$

$$\begin{aligned} avgT_{turnaround} &= \\ &= ((26 - 0) + (28 - 0) + (30 - 0))/3 = 28 \end{aligned}$$

# The length of the time slice is critical

- The shorter the time slice ...
  - ... the better the response time.
  - ... the higher the relative cost of context switching.
- The longer time slice ...
  - ... the less noticeable the increase in the cost of switching.
  - ... the worse the response time.

# Relaxed Workload Assumptions

- ~~1. All jobs run for the same amount of time.~~
- ~~2. All jobs arrive at the same time.~~
- ~~3. All jobs only use the CPU (i.e. no I/O).~~
4. Run-time of each job is known.

# Incorporating I/O

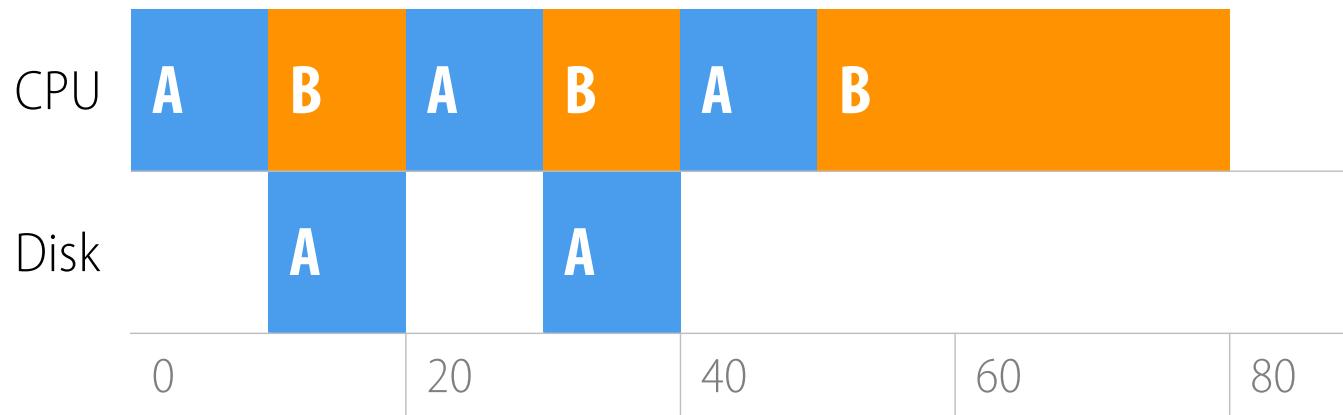
- Let's relax assumption 3: programs may perform I/O
- Example:
  - Jobs A and B need 50 units of CPU time each.
  - A runs for 10 units and then issues an I/O request
    - Each I/O operation takes 10 time units
  - B simply uses the CPU for 50 time units and performs no I/O
  - The scheduler runs A first, then B.

# What happens when the scheduler isn't I/O aware?



- Resources may be poorly used.
- Job A holds on to CPU even while blocked waiting for disk.

# Making the scheduler I/O aware



- Treat Job A as 3 separate CPU bursts.
- When Job A completes I/O, another Job A becomes ready.
- Each CPU burst is shorter than Job B, so even with SCTF, Job A preempts Job B.

# Making the scheduler I/O aware

- When a job initiates an I/O request...
  - ... the job is blocked waiting for I/O completion.
  - ... the scheduler should schedule another job on the CPU.
- When the I/O completes...
  - ... an interrupt is raised.
  - ... the OS moves the process from blocked back to ready state.

# Can we further relax our Workload Assumptions?

- ~~1. All jobs run for the same amount of time.~~
- ~~2. All jobs arrive at the same time.~~
- ~~3. All jobs only use the CPU (i.e. no I/O).~~
- ~~4. Run time of each job is known.~~

Well... that's a theme for next week.