

T10


Memory Virtualization Segmentation

Referência principal

Ch.16 of Operating Systems: Three Easy Pieces by Remzi and Andrea Arpaci-Dusseau (pages.cs.wisc.edu/~remzi/OSTEP/)

Discutido em classe em 29 de agosto de 2018

Inefficiency of the Base and Bound Approach



- Big chunk of “free” space
- “Free” space takes up physical memory.
- Hard to run when an address space does not fit into physical memory

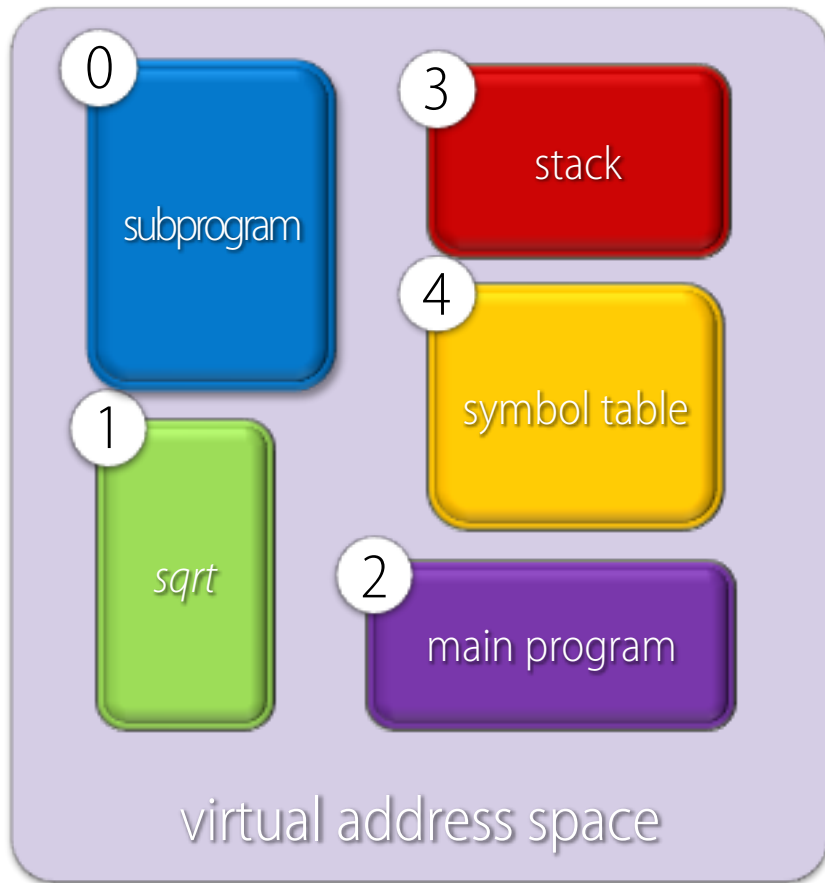
Segmentation

- Memory-management scheme that supports user view of memory.
 - A program is a collection of segments.
- A segment is a virtual unit such as:
 - main program
 - procedure
 - function
 - method
 - object
 - local variables
 - global variables
 - common block
 - stack
 - symbol table

Segmentation


- A segment is just a contiguous portion of the address space of particular length.
 - Logically-different segments: code, stack, heap
- Each segment can be placed in a different part of physical memory.
 - Base and bound are defined for each segment.

Logical View of Segmentation



	base	bound
0	1400	1000
1	6300	400
2	43000	400
3	3200	1100
4	4700	1000

segment table



Segmentation

- Segmentation provides a virtual view of main memory.
- The segments of a program may have different lengths.
- There is a maximum segment length.
- Since segments are not equal, segmentation resembles dynamic partitioning.

Segmentation

- Segment is a contiguous region of virtual memory
- Segment can be located anywhere in physical memory
 - Each segment has: start, length, access permission
- Processes can share segments
 - Same start, length, same/different access permissions

Segmentation Architecture

- Virtual address consists of a pair:
 <segment-number, offset>
- Each process has a segment table (in hardware)
 - Segment number → entry in table
- A segment table maps a two-dimensional virtual address onto a one-dimensional physical address.
 - Each entry in a segment table has at least two fields
 - **Base**: the segment's starting physical address in memory.
 - **Bound**: the length of the segment.

Segmentation Architecture

- Protection
 - With each entry in segment table associate
 - validation bit (if zero \Rightarrow illegal segment)
 - privilege bits (read/write/execute)
 - Protection bits associated with segments
 - Code sharing occurs at segment level.
- Since segments vary in length, memory allocation is a dynamic storage allocation problem.

Segmentation Architecture

Two hardware registers are used to keep track of the segment table


- The **Segment-Table Base Register (STBR)** points to the segment table's location in memory.
- The **Segment-Table Length Register (STLR)** indicates the number of segments used by a program
 - Segment number s is legal if $s < STLR$.

Fine-Grained and Coarse-Grained Segmentation

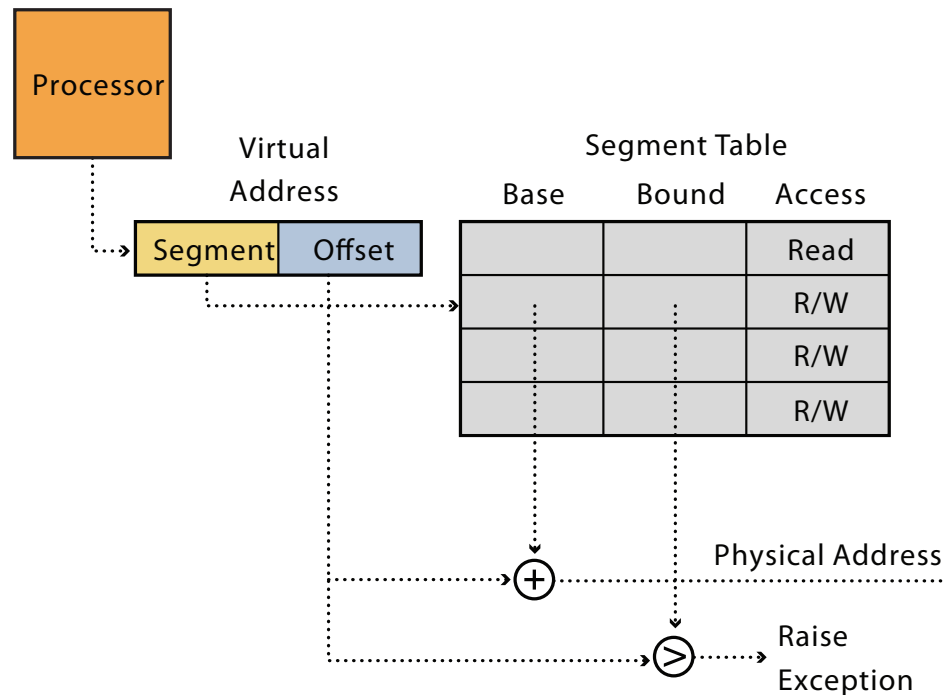
- **Coarse-Grained** means segmentation in small numbers.
 - e.g., code, heap, stack.
- **Fine-Grained** segmentation allows a more flexible address space.
 - To support many segments, hardware support with a **segment table** is required.

Segmentation

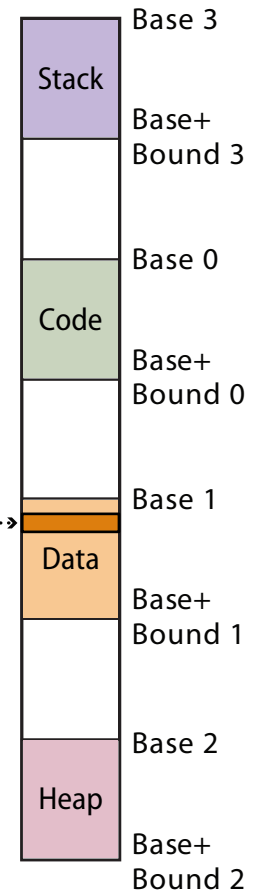
Processor's View



Implementation



Physical Memory



Segmentation

■ Pros?

- Can share code/data segments between processes
- Can protect code segment from being overwritten
- Can transparently grow stack/heap as needed
- Can detect if need to copy-on-write or zero-on-reference

■ Cons?


- Complex memory management
 - Need to find chunk of a particular size
- May need to rearrange memory from time to time to make room for new segment or growing segment
 - External fragmentation: wasted space between chunks

UNIX fork and Copy on Write


- UNIX fork
 - Makes a complete copy of a process
- Segments allow a more efficient implementation
 - Copy segment table into child
 - Mark parent and child segments as read-only
 - Start child process; return to parent
 - If child or parent writes to a segment (ex: stack, heap)
 - trap into kernel
 - make a copy of the segment and resume

Processor's View

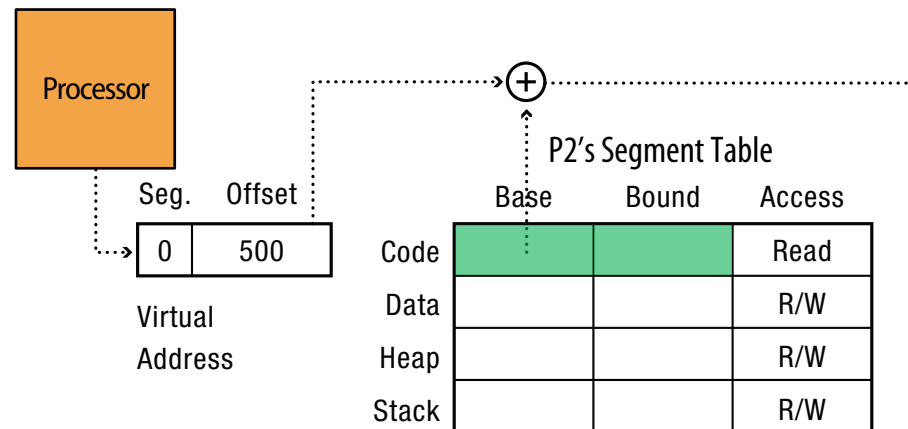
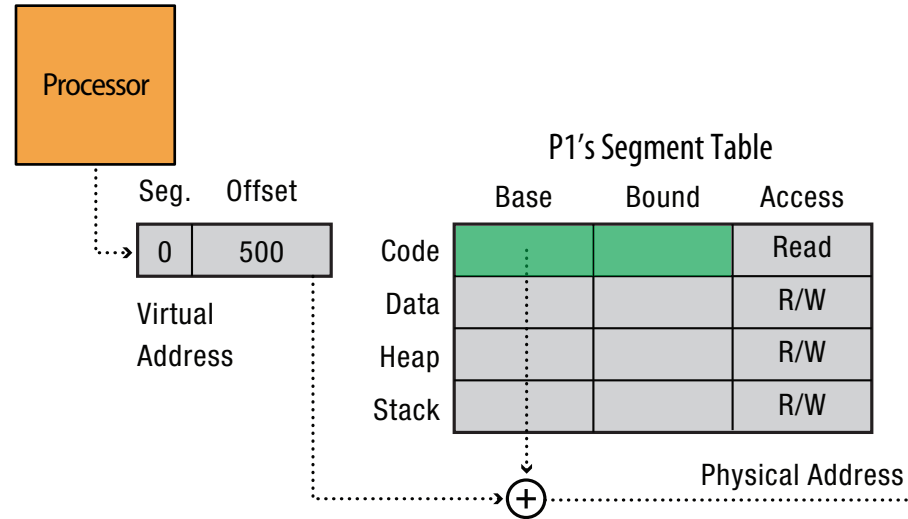
Process 1's View




Process 2's View



Implementation



Physical Memory





Zero-on-Reference

- How much physical memory is needed for the stack or heap?
 - Only what is currently in use
- When program uses memory beyond end of stack
 - Segmentation fault into OS kernel
 - Kernel allocates some memory
 - How much?
 - Clears the allocated memory
 - Avoid accidentally leaking information!
 - Modify segment table
 - Resume process

A point to ponder

With segmentation, what is saved/restored on a process context switch?

Placing Segments In Physical Memory

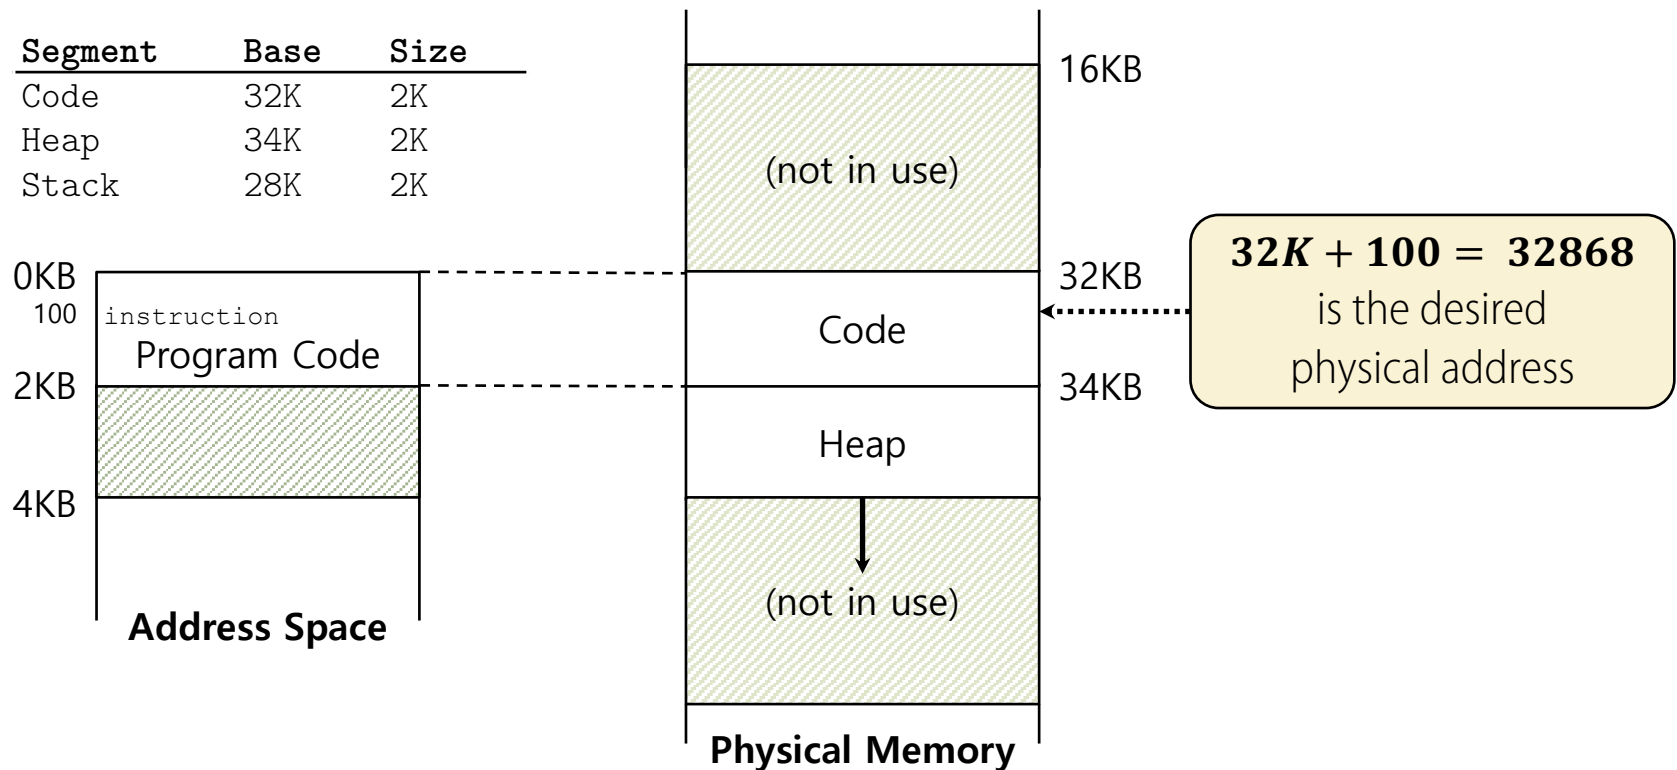


Segment	Base	Size
Code	32K	2K
Heap	34K	2K
Stack	28K	2K

Example

Address Translation on Segmentation

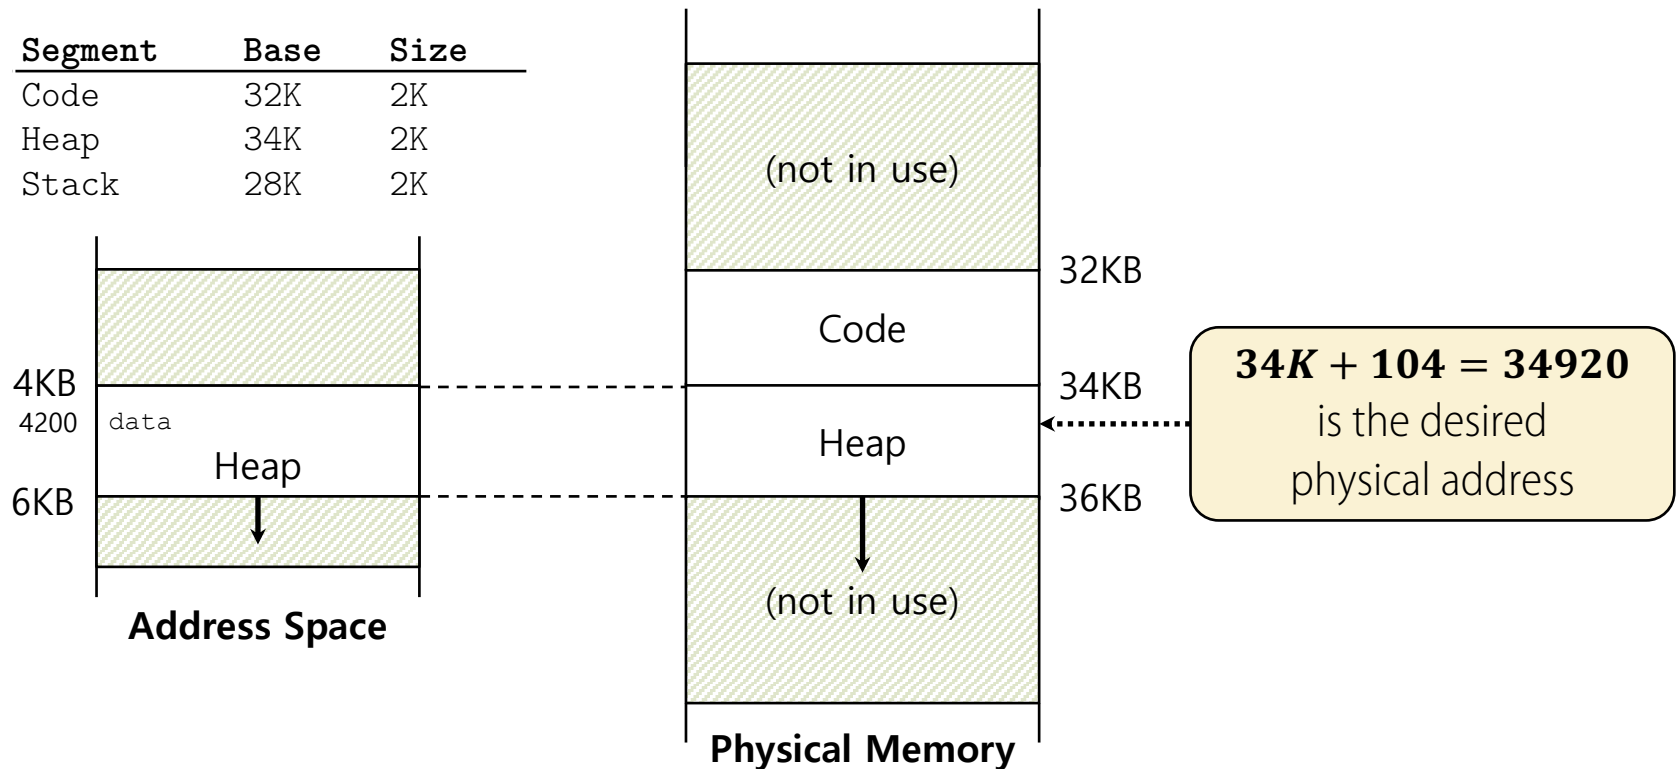
- The code segment starts at virtual address **0** in address space.
- The offset of virtual address **100** is **100**.
- $physical\ address = base + offset$



Example

Address Translation on Segmentation


- The heap segment starts at virtual address **4096** in address space.
- Thus, the offset of virtual address **4200** is $4200 - 4096 = 104$.
- *physical address = base + offset.*



Segmentation Fault or Violation


- If an **illegal address** such as 7KB which is beyond the end of heap is referenced, the OS incurs a **segmentation fault**.
 - The hardware detects that address is **out of bounds**.

Segment	Base	Size
Code	32K	2K
Heap	34K	2K
Stack	28K	2K




OS support

Which Segment Are We Referring To?




- The hardware uses segment registers during address translation. **How does it know the offset into a segment, and to which segment an address refers?**
- One possible approach is to chop up the address space into segments based on the **top few bits** of the virtual address.
 - For example, if the OS requires three segments, the top two bits of the virtual address could be used to represent the segment number.
 - Assuming an address space of 16KB, the remaining 12 bits will be used to represent the offset into the segment, which leaves us with a maximum segment size of 4KB.



Example

Referring to a Segment

- Let us assume that we want to know what segment and offset a virtual address **4200** refers to.
 - Since $4200_{10} = 01000001101000_2$ that address would be read as



- Assuming that our segments are numbered as shown on the table, virtual address **4200**₁₀ would refer to an offset of **104**₁₀ into the **heap** (segment #1).

Segment	bits
Code	00
Heap	01
Stack	10
(unused)	11

Referring to a Segment


- Assuming that **base** and **bounds** were arrays indexed by segment number, to translate a virtual address into a physical one, the hardware would be doing something like...

```
1  #define SEG_MASK 0x3000    // (1100000000000000)
2  #define SEG_SHIFT 12
3  #define OFFSET_MASK 0xFFF  // (0011111111111111)
4
5  // get top 2 bits of 14-bit VA
6  Segment = (VirtualAddress & SEG_MASK) >> SEG_SHIFT;
7  // get offset
8  Offset = VirtualAddress & OFFSET_MASK;
9  if (Offset >= Bounds[Segment])
10     RaiseException(PROTECTION_FAULT);
11 else {
12     PhysicalAddress = Base[Segment] + Offset;
13     Register = AccessMemory(PhysicalAddress);
14 }
```


OS support

Referring to the Stack Segment

- Remember that **the stack grows backwards**.
- **Extra hardware support** is needed.
 - The hardware must check which way a segment grows.
 - Let us use a bit for that: 1 = positive direction, 0 = negative direction




Segment Registers with Support for Negative-Growth

Segment	Base	Size	Positive-Growth?
Code	32K	2K	1
Heap	34K	2K	1
Stack	28K	2K	0

Example

Referring to an Address within the Stack

- Let us assume that we want to know what segment and offset a virtual address **15K** refers to.
 - Since $15K = 11110000000000_2$ that address would be read as




- Assuming that our segments are numbered as shown on the table, virtual address 4200_{10} would refer to an offset of 104_{10} into the **heap** (segment #1).

Segment	bits
Code	00
Heap	01
Stack	10
(unused)	11

Example

Referring to an Address within the Stack

- Let us assume that we want to know what segment and offset a virtual address **15K** refers to.
 - Since **15K** = **11110000000000**₂ that address would be read as



- Since the stack segment grows backwards the **3K** offset must be subtracted from the maximum segment size (**4K**) to obtain the correct negative offset.

Segment	Base	Size	Positive-Growth?
Code	32K	2K	1
Heap	34K	2K	1
Stack	28K	2K	0

- 3K** − **4K** = **−1K** which will now be added to the stack base address (**28K**) to generate the correct physical address (**27K**).
- The offset is also checked against bounds by comparing its absolute value with the segment's size.

Sharing

- Segments can also be **shared among address spaces**.
 - **Code sharing**, for example, is common and still used in systems today.
- Segment sharing requires extra hardware support in the form of **protection bits**.
 - A few more bits are added to the segment registers or to the segment table to indicate permission to read, write and/or execute.

Segment Registers with Support for Negative-Growth and Sharing

Segment	Base	Size	Positive-Growth?	Protection
Code	32K	2K	1	Read-Execute
Heap	34K	2K	1	Read-Write
Stack	28K	2K	0	Read-Write

Fragmentation

■ External Fragmentation


- Holes of free space in physical memory that make it difficult to allocate new segments.
 - There are 24KB free, but not in one contiguous segment.
 - The OS cannot satisfy a 20KB request.

■ Compaction

- Rearranging the existing segments in physical memory, but this is expensive.
 - Stop running processes.
 - Copy data to somewhere.
 - Change segment base register value.

Memory Compaction

BEFORE



AFTER

