



Memory Virtualization Paging

Referência principal

Ch.18 of *Operating Systems: Three Easy Pieces* by Remzi and Andrea Arpaci-Dusseau (pages.cs.wisc.edu/~remzi/OSTEP/)

Discutido em classe em 03 de setembro de 2018

Partitioning Issues

- We have seen that
 - With fixed partitioning, any program, no matter how small, occupies an entire partition.
 - This sort of memory waste is called **internal fragmentation**.
 - With variable partitioning, due to the process allocation policy, holes of various sizes may appear and be scattered throughout memory.
 - This sort of memory waste is called **external fragmentation**.

Noncontiguous allocation

- One way to fight external fragmentation is to make the physical address space of a process noncontiguous.
 - The idea is to make all the available physical memory accessible to any running process.
- We have already studied segmentation, an early noncontiguous allocation concept, which reminds us of the pros and cons of variable partitioning.
- Paging is another early noncontiguous allocation concept, which takes a different approach.
 - It was designed at the University of Manchester and implemented in the Atlas Computer, which was commissioned in 1962.

The Paging Approach

- The idea is to use fixed-sized units instead of variable-sized logical segments.
- Split virtual memory into blocks of the same size (**pages**).
- Split physical memory into fixed-sized blocks (**frames**).
 - Frame size is 2^x bytes (today usually $9 \leq x \leq 13$ but it can be as high as **31**).
- Keep track of all free frames.
- To run a program with n pages, load it into n free frames.
- Set up a **page table** to map virtual pages to the associated physical frames.

Points to ponder

- How can we virtualize memory with pages, so as to avoid the problems of segmentation?
- What are the basic techniques?
- How do we make those techniques work well, with minimal space and time overheads?
- Is this strategy prone to...
 - Internal fragmentation? Why?
 - External fragmentation? Why?

Advantages Of Paging

■ Flexibility

- It supports the abstraction of an address space effectively.
 - E.g. there will be no assumptions about how heap and stack grow and are used.


■ Simplicity

- Managing the page table requires only two hardware registers
 - Page-table base register (**PTBR**) points to page table.
 - Page-table length register (**PTLR**) indicates the size of the page table.
- It is easy to manage the free-space.
 - The pages in the address space and the frames in the physical memory are the same size.
 - It is easy to allocate and keep a free list, e.g. by using a bitmap like 0011111100000001100, where each bit represents one physical page frame.

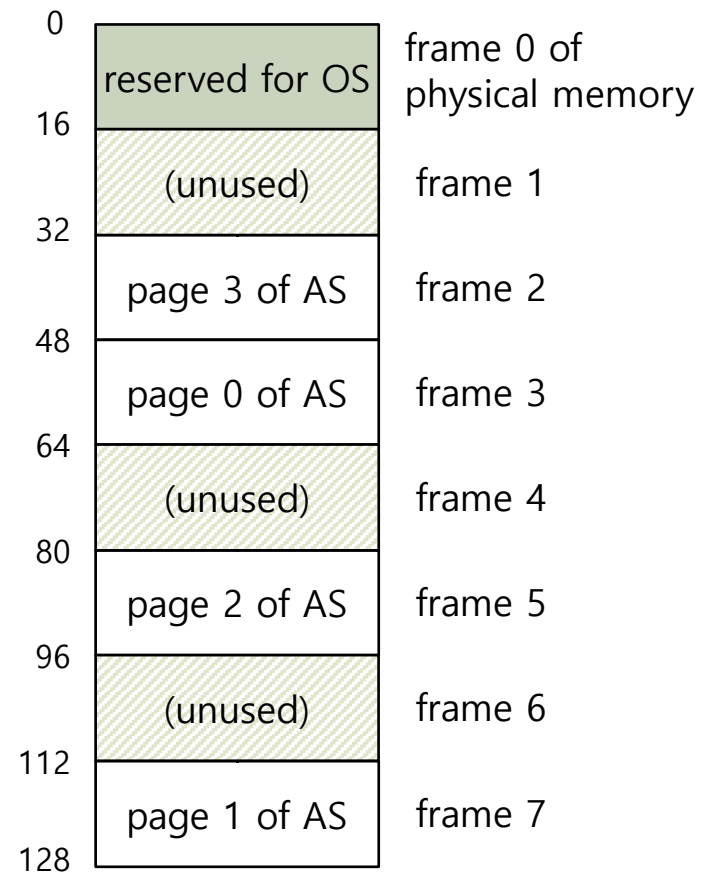
Example

A Simple Paging System

- Consider
 - A tiny 64B address space with four 16B pages
 - A small 128B physical memory with eight 16B page frames.




A Simple 64B Address Space



The Address Space Placed In Physical Memory

Example


A Simple Paging System




virtual
memory

page number	
0	3
1	7
2	5
3	2

page table



Paged Translation (Abstract)



Address Translation Scheme

- Virtual address generated by CPU is divided into two parts


Virtual Page Number (*VPN*)

Used as an index into a page table which contains the base address of each page in physical memory.

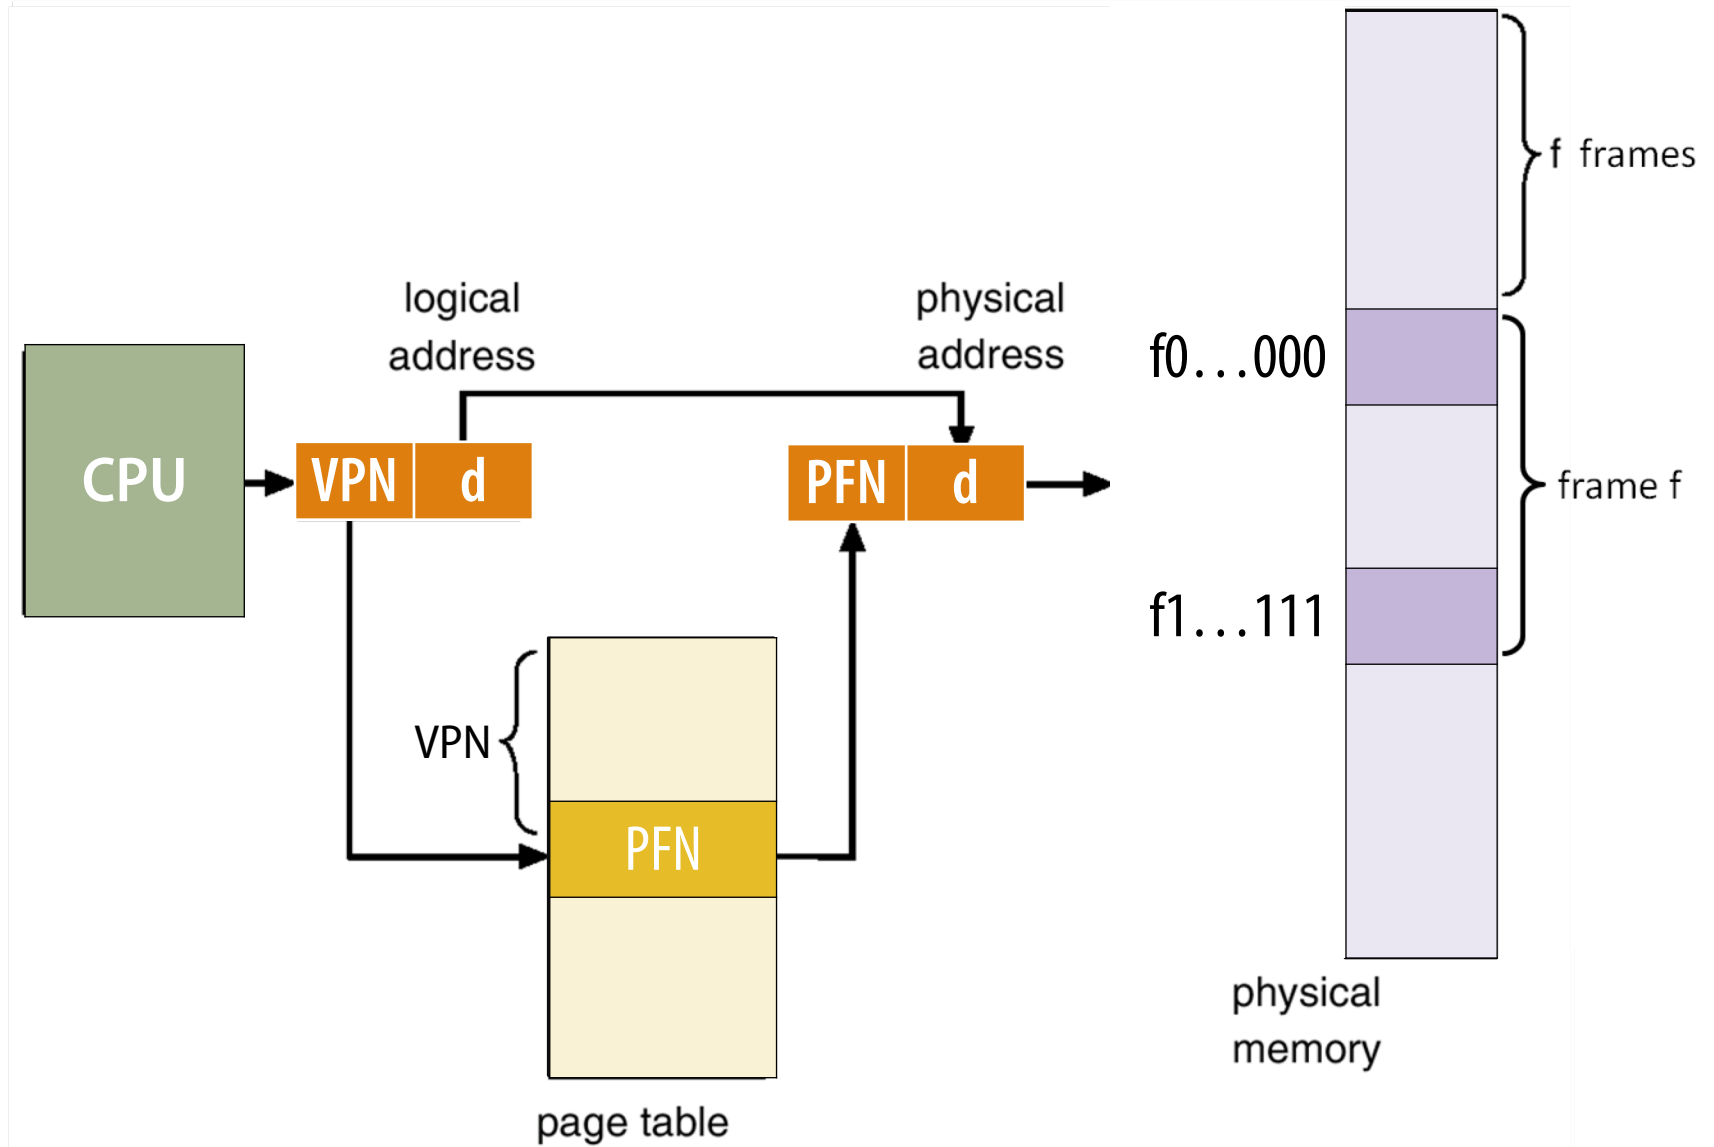
Page offset (*d*)

Combined with base address to define the physical memory address that is sent to the memory unit.

- In our simple paging example, 6 bits will be required to cover the address space.
 - There are 4 pages, so the VPN will take 2 bits
 - The size of each page is 16B, so the offset will require 4 bits.




Address Translation Scheme



Example

Address Translation

- Consider the 64B address space in the simple paging scheme example.
- Take virtual address 21_{10} and let us decompose it
 - $21_{10} = 010101_2$, which means
 - VPN = 1 and
 - Offset = 5B




Example

Address Translation

- Assume now that, in our simple system, Virtual Pages are mapped onto Physical Frames according to the table on the left.
- Virtual address 21_{10} would be mapped onto memory address 117_{10} as shown by the diagram on the right.

VPN	PFN
0	3
1	7
2	5
3	2



Paging Example (4B pages, 32B mem)

VA ₁₀	VA ₄	Value
0	00	A
1	01	B
2	02	C
3	03	D
4	10	E
5	11	F
6	12	G
7	13	H
8	20	I
9	21	J
10	22	K
11	23	L
12	30	M
13	31	N
14	32	O
15	33	P


Page Table	
0	5
1	6
2	3
3	1
4	
5	
6	
7	

PA ₄	PA ₁₀	Value
00	0	
01	1	
02	2	
03	3	
10	4	M
11	5	N
12	6	O
13	7	P
20	8	
21	9	
22	10	
23	11	
30	12	I
31	13	J
32	14	K
33	15	L

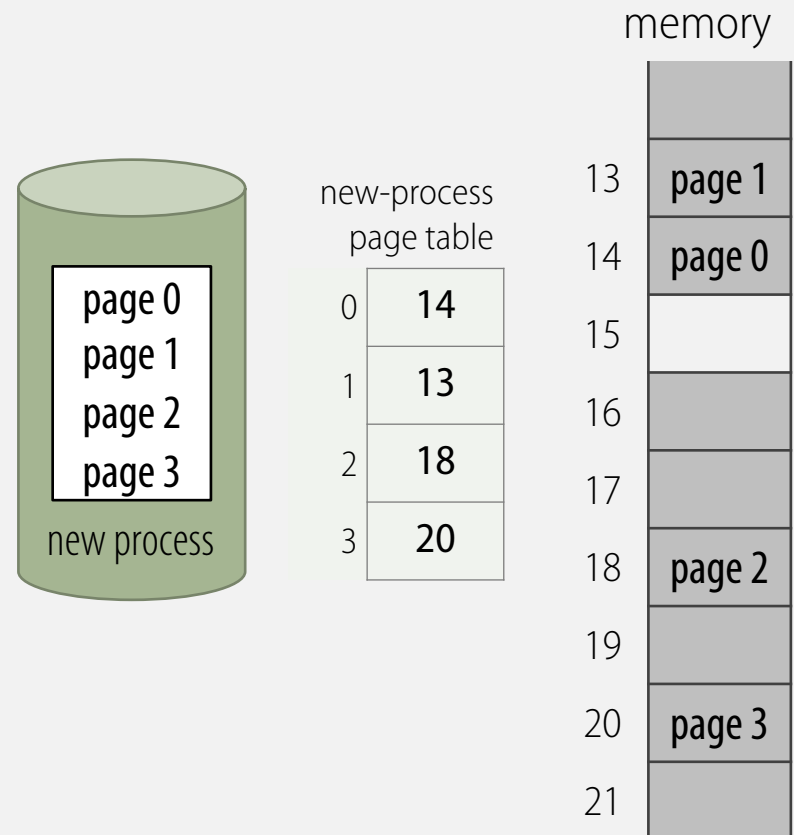
PA ₄	PA ₁₀	Value
100	16	
101	17	
102	18	
103	19	
110	20	A
111	21	B
112	22	C
113	23	D
120	24	E
121	25	F
122	26	G
123	27	H
130	28	
131	29	
132	30	
133	31	

Free Frame Allocation

BEFORE



AFTER



Points to ponder...

- How big can page tables be?
- Where are page tables stored?
- What are the typical contents of a page table?
- Does paging make the system (too) slow?

How Big Are Page Tables?


- Page tables can get awfully large, much larger than segment tables or base/bound registers that we have already discussed.
 - For example, take a 32-bit address space with 4KB pages.
 - Generating every possible address within a page requires **12** bits ($2^{12} = 4K$).
 - Since there are 32 bits in the address space and we have used 12, there will be 20 bits left, which means 2^{20} pages.
 - Assuming that each entry in a page table takes 4B, this page table will occupy $2^{20} \times 4B = 4MB$.
 - And there will be one of these tables for each running process...

Where Are Page Tables Stored?

- Because page tables are so big, they cannot fit into any special on-chip hardware.
- Instead, they are kept in the kernel physical memory.

What Is In The Page Table?

- The page table is just a **data structure** that is used to map a virtual address to a physical address, and in its simplest form is just an array.
- The OS **indexes** the PT by VPN to get the Physical Frame Number and some flags, e.g.
 - **Valid (V)**: whether the frame is valid.
 - **Read/Write (R/W)**: whether the page can be read from or written to
 - **Present (P)**: whether the page is in physical memory or on disk
 - **Dirty (D)**: whether the page has been modified since it was read
 - **Accessed (A)**: whether the page has been accessed



Example: An x86 Page Table Entry (PTE)

Paging: Also Too Slow

- We have seen that in order to translate a virtual address, we need to access the Page Table to get the corresponding Physical Frame Number.
- Since the Page Table is stored in memory, this implies that, for every memory reference, paging requires the OS to perform **one extra memory reference**.
- An abstract protocol for what happens on each memory reference is shown on the next slide.

Accessing Memory With Paging

```
1  // Extract the VPN from the virtual address
2  VPN = (VirtualAddress & VPN_MASK) >> SHIFT;
3
4  // Form the address of the page-table entry (PTE)
5  PTEAddr = PTBR + (VPN * sizeof(PTE));
6
7  // Fetch the PTE
8  PTE = AccessMemory(PTEAddr);
9
10 // Check if process can access the page
11 if (PTE.Valid == False)
12     RaiseException(SEGMENTATION_FAULT);
13 else if (CanAccess(PTE.ProtectBits) == False)
14     RaiseException(PROTECTION_FAULT);
15 else {
16     // Access is OK: form physical address and fetch it
17     offset = VirtualAddress & OFFSET_MASK;
18     PhysAddr = (PTE.PFN << PFN_SHIFT) | offset;
19     Register = AccessMemory(PhysAddr);
20 }
```

Points to Ponder...

With paging,
what is saved or
restored on a
process context
switch?

What if the page
size is very small?

What if the page
size is very large?

Paging and Copy on Write

■ Can we share memory between processes?

- Set entries in both page tables to point to same page frames
- Need core map of page frames to track which processes are pointing to which page frames (e.g., reference count)

- UNIX fork with copy on write
 - Copy page table of parent into child process
 - Mark all pages (in new and old page tables) as read-only
 - Trap into kernel on write (in child or parent)
 - Copy page
 - Mark both as writeable
 - Resume execution

Fill On Demand

- **Can I start running a program before its code is in physical memory?**
 - Set all page table entries to invalid
 - When a page is referenced for first time, kernel trap
 - Kernel brings page in from disk
 - Resume execution
 - Remaining pages can be transferred in the background while program is running

Sparse Address Spaces

- Might want many separate dynamic modules
 - Per-processor heaps
 - Per-thread stacks
 - Memory-mapped files
 - Dynamically linked libraries
- **What if virtual address space is large?**
 - 32-bits, 4KB pages → 500K page table entries
 - 64-bits, 4KB pages → 4 quadrillion page table entries