



Semaphores

Referência principal

Ch.31 of *Operating Systems: Three Easy Pieces* by Remzi and Andrea Arpaci-Dusseau (pages.cs.wisc.edu/~remzi/OSTEP/)

Discutido em classe em 15 de outubro de 2018

Semaphores: A Definition

- A semaphore is a synchronization object with an integer value that, in POSIX, can be manipulated by two routines: `sem_wait()` and `sem_post()`.
- Before being used, a semaphore must be initialized and this initial value will determine its behavior, as we will see next.

```
1. #include "semaphore.h"
2. sem_t s;
3. sem_init(&s, 0, 1);
```

On linux, "semaphore.h" defaults to <semaphore.h>.
On macOS it loads a wrapper that simulates the POSIX API.

- Here, a semaphore `s` is created and initialized to `1` (the third argument of the call).
- The second argument is zero to indicate that the semaphore will be shared among threads in the same process.
 - A semaphore can also be shared among several processes, and this is what other values of this second argument would indicate.

Semantics of `sem_wait()`

```
1. int sem_wait(sem_t *s) {  
2.     decrement the value of semaphore s by one;  
3.     if value of semaphore s is negative  
4.         wait on s;  
5. }
```

- `sem_wait()` is implemented atomically.
- It decrements the value of the semaphore and returns immediately if the result is non-negative.
- Otherwise, it will cause the caller to sleep until awakened by a subsequent `sem_post()` call from another thread.

Semantics of `sem_post()`

```
1.  int sem_post(sem_t *s) {  
2.    increment the value of semaphore s by one;  
3.    if there are one or more threads waiting on s  
4.      wake one up  
5.  }
```

- `sem_post()` is also implemented atomically.
- It increments the semaphore and returns immediately if its associated waiting list is empty.
- Otherwise, it will also awake a waiting thread before returning.

Binary Semaphores (Locks)

- A binary semaphore functions just like a lock

```
1. sem_t bs;  
2. sem_init(&bs, 0, X) ; // initialize semaphore to X. What should X be?  
  
3. sem_wait(&bs);  
4. // critical section here  
5. sem_post(&bs);
```

- Look back at the definition of `sem_wait()` and `sem_post()` and choose the right value for `X`.
- To discuss the functioning of the binary semaphore, let us examine a scenario with two threads.

Thread trace: a single thread uses the semaphore

- In this case, there are two threads, but Thread 0 runs without interruption.

Value of Semaphore	Thread 0	Thread 1
1		
1	call <code>sem_wait()</code>	
0	<code>sem_wait()</code> returns	
0	(critical section)	
0	call <code>sem_post()</code>	
1	<code>sem_post()</code> returns	

Thread trace: two threads use the semaphore

- In this case, there are two threads, which compete for the semaphore.
- Thread 0 is interrupted twice by the system.
- Experiment with scenarios of your choice to make sure you understand the model.

Value	Thread 0		Thread 1	
	state	action	state	action
1	Running		Ready	
1	Running	call <code>sem_wait()</code>	Ready	
0	Running	<code>sem_wait()</code> returns	Ready	
0	Running	(crit sect: begin)	Ready	
0	Ready	<i>Interrupt: switch → T1</i>	Running	
0	Ready		Running	call <code>sem_wait()</code>
-1	Ready		Running	decrement <code>sem</code>
-1	Ready		Sleeping	(<code>sem < 0</code>) → asleep
-1	Running		Sleeping	<i>Switch → T0</i>
-1	Running	(crit sect: end)	Sleeping	
-1	Running	call <code>sem_post()</code>	Sleeping	
0	Running	increment <code>sem</code>	Sleeping	
0	Running	wake (T1)	Ready	
0	Running	<code>sem_post()</code> returns	Ready	
0	Ready	<i>Interrupt: switch → T1</i>	Running	
0	Ready		Running	<code>sem_wait()</code> returns
0	Ready		Running	(critical section)
0	Ready		Running	call <code>sem_post()</code>
1	Ready		Running	<code>sem_post()</code> returns

Semaphore as Condition Variable: Parent Waiting for Child

- We have done this before: a thread creates another thread and then waits until it finishes.
- What should the value of X in line 8 be to achieve the desired functionality?
- Again, let us examine two thread traces to get acquainted with the algorithm.

```
1.  sem_t s;  
  
2.  void *child(void *arg) {  
3.      printf("child\n");  
4.      sem_post(&s); // child is done  
5.      return NULL;  
6.  }  
  
7.  int main(int argc, char *argv[]) {  
8.      sem_init(&s, 0, X); // what should X be?  
9.      printf("parent: begin\n");  
10.     pthread_t c;  
11.     mythread_create(&c, NULL, child, NULL);  
12.     sem_wait(&s); // wait here for child  
13.     printf("parent: end\n");  
14.     return 0;  
15. }
```


Thread trace: parent waiting for child (case 1)

- In this case, we assume that the parent thread is not interrupted after having created the child.
- Thus, the parent calls `sem_wait()` before the child calls `sem_post()`.

Value	Parent	State	Child	State
0	create Child	Running	<i>Child exists; is runnable</i>	Ready
0	call <code>sem_wait ()</code>	Running		Ready
-1	decrement <code>sem</code>	Running		Ready
-1	$(sem < 0) \rightarrow$ asleep	Sleeping		Ready
-1	<i>Switch \rightarrow Child</i>	Sleeping	child runs	Running
-1		Sleeping	call <code>sem_post ()</code>	Running
0		Sleeping	increment <code>sem</code>	Running
0		Ready	wake Parent	Running
0		Ready	<code>sem_post()</code> returns	Running
0		Ready	<i>Interrupt: Switch \rightarrow Parent</i>	Ready
0	<code>sem.wait ()</code> returns	Running		Ready

Thread trace: parent waiting for child (case 2)

- In this case, we assume that the parent thread is interrupted just after having created the child.
- Thus, the child calls `sem_post()` before the parent calls `sem_wait()`.

Value	Parent	State	Child	State
0	create Child	Running	<i>Child exists; is runnable</i>	Ready
0	<i>Interrupt: Switch → Child</i>	Ready	child runs	Running
0		Ready	call <code>sem_post()</code>	Running
1		Ready	increment <code>sem</code>	Running
1		Ready	wake nobody	Running
1		Ready	<code>sem_post()</code> returns	Running
1	parent runs	Running	<i>Interrupt: Switch → Parent</i>	Ready
1	call <code>sem_wait()</code>	Running		Ready
0	decrement <code>sem</code>	Running		Ready
0	<code>(sem > 0) → awake</code>	Running		Ready

Solving the Producer / Consumer Problem

- In this case we will use what is called a *counting semaphore*.
- We will introduce two semaphores, **empty** and **full**, that the threads will use to indicate that a buffer entry has been emptied or filled, respectively.

```
1. sem_t empty;
2. sem_t full;

3. int main(int argc, char *argv[]) {
4.     // ...
5.     sem_init(&empty, 0, MAX); // MAX buffers are empty to begin with...
6.     sem_init(&full, 0, 0);    // ... and 0 are full
7.     // ...
```

The `put()` and `get()` routines

- These are essentially the same routines we designed for our solution using condition variables.

```
1. int buffer[MAX];
2. int fill = 0;
3. int use = 0;

4. void put(int value) {
5.     buffer[fill] = value;    // f1
6.     fill = (fill + 1) % MAX; // f2
7. }

8. int get(void) {
9.     int tmp = buffer[use];    // g1
10.    use = (use + 1) % MAX;    // g2
11.    return tmp;
12. }
```

Adding **full** and **empty** semaphores

```
1. void *producer(void *arg) {
2.     for (int i = 0; i < loops; i++) {
3.         sem_wait(&empty) ;           // p1
4.         put(i);                       // p2
5.         sem_post(&full);              // p3
6.     }
7. }

8. void *consumer(void *arg) {
9.     for (int i = 0; i < loops; i++) {
10.        sem_wait(&full);                // c1
11.        int tmp = get();                // c2
12.        sem_post(&empty);               // c3
13.        printf("%d\n", tmp);
14.    }
15. }
```

- Assume $MAX = 1$, one producer and one consumer.
 - Does it work?
- Assume $MAX = 10$, one producer and one consumer.
 - Does it work?
- Assume more than one producers and consumers.
 - Does it work?

Trying to implement mutual exclusion

- When there are multiple producers or consumers a race condition arises in `put()` and `get()`.
- We can try to prevent it using a binary semaphore to implement mutual exclusion among those calls.
- We'll create it in `main()` and later use it in `producer()` and `consumer()`.

```
1.  sem_t empty;  
2.  sem_t full;  
3.  sem_t mutex;  
  
4.  int main(int argc, char *argv[]) {  
5.      // ...  
6.      sem_init(&empty, 0, MAX); // MAX buffers are empty to begin with...  
7.      sem_init(&full, 0, 0);    // ... and 0 are full  
8.      sem_init(&mutex, 0, 1);  // mutex = 1 because it is a lock  
9.      // ...  
10. }
```

Adding **mutex** to **producer()** and **consumer()**

```
1. void *producer(void *arg) {  
2.     for (int i = 0; i < loops; i++) {  
3.         sem_wait(&mutex);  
4.         sem_wait(&empty);  
5.         put(i);  
6.         sem_post(&full);  
7.         sem_post(&mutex);  
8.     }  
9. }
```

```
1. void *consumer(void *arg) {  
2.     for (int i = 0; i < loops; i++) {  
3.         sem_wait(&mutex);  
4.         sem_wait(&full);  
5.         int tmp = get( );  
6.         sem_post(&empty);  
7.         sem_post(&mutex);  
8.         printf("%d\n", tmp);  
9.     }  
10. }
```

- Here we added **mutex** to **producer()** and **consumer()** to prevent the race condition on the buffer.
- Is our solution correct now?

How to avoid deadlock

- Our proposed scheme does not work because `producer()` and `consumer()` may get trapped, each one waiting for a condition that only the other could provide.
- This is a classical problem known as **deadlock**, which we will study in detail later.
- In this case, the problem can be solved by reducing the scope of the mutual exclusion between `producer()` and `consumer()`, as shown in the next slide.
- The result is a simple and working bounded buffer, a commonly-used pattern in multithreaded programs.

Avoiding deadlock

```
1. void *producer(void *arg) {  
2.     for (int i = 0; i < loops; i++) {  
3.         sem_wait(&empty);  
4.         sem_wait(&mutex);  
5.         put(i);  
6.         sem_post(&mutex);  
7.         sem_post(&full);  
8.     }  
9. }
```

- Here we reduce the scope of `mutex` in `producer()` and `consumer()` to prevent the deadlock.
- Is our solution correct now?
- Why do producer and consumer wait and signal different semaphores?

```
1. void *consumer(void *arg) {  
2.     for (int i = 0; i < loops; i++) {  
3.         sem_wait(&full);  
4.         sem_wait(&mutex);  
5.         int tmp = get( );  
6.         sem_post(&mutex);  
7.         sem_post(&empty);  
8.         printf("%d\n", tmp);  
9.     }  
10. }
```

- Is order of waits important?
- Is order of signals important?
- What if we have 2 producers or 2 consumers? Do we need to change anything?
- Can we use semaphores for FIFO ordering?

Reader-Writer Locks

- In some cases we need a more flexible primitive that takes into account that different accesses to a data structure might require different kinds of locking.
- For example, imagine a number of concurrent list operations, including inserts and simple lookups.
 - Inserts change the state of the list and thus must be protected by e.g. a critical section.
 - On the other hand, lookups simply read the list and, as long as there is no simultaneous insert, many of them may proceed concurrently.
- The special type of lock that will support this type of operation is known as a **reader-writer lock**.

A simple reader-writer lock

```
1. typedef struct _rwlock_t {
2.     sem_t lock;        // binary semaphore (basic lock)
3.     sem_t writelock;    // used to allow ONE writer or MANY readers
4.     int readers;        // count of readers reading in critical section
5. } rwlock_t;

6. void rwlock_init(rwlock_t *rw) {
7.     rw->readers = 0;
8.     sem_init(&rw->lock, 0, 1) ;
9.     sem_init(&rw->writelock, 0, 1);
10. }
```

A simple reader-writer lock

```
1. void rwlock_acquire_readlock(rwlock_t *rw) {
2.     sem_wait(&rw->lock);
3.     rw->readers++;
4.     if (rw->readers == 1)
5.         sem_wait(&rw->writelock); // first reader acquires writelock
6.     sem_post(&rw->lock);
7. }

8. void rwlock_release_readlock(rwlock_t *rw) {
9.     sem_wait(&rw->lock);
10.    rw->readers--;
11.    if (rw->readers == 0)
12.        sem_post(&rw->writelock); // last reader releases writelock
13.    sem_post(&rw->lock);
14. }
```


A simple reader-writer lock

```
1.  void rwlock_acquire_writelock(rwlock_t *rw) {  
2.      sem_wait(&rw->writelock) ;  
3.  }  
  
4.  void rwlock_release_writelock(rwlock_t *rw) {  
5.      sem_post(&rw->writelock) ;  
6.  }
```

- Examine the assumptions and the design of this solution.
 - Is it correct? Is it fair? Is it efficient?
- Can readers harm writers somehow? Can writers harm readers?

The Dining Philosophers

- Assume there are five “philosophers” sitting around a table.
- Between each pair of philosophers is a single hashi (and thus, five total).
- The philosophers each have times where they think, and don't need any hashis, and times when they eat.
- In order to eat, a philosopher needs two hashis, both the one on their left and the one on their right.
- The contention for these hashis, and the synchronization problems that ensue, are what makes this problem worth of being studied in concurrent programming.



Designing a solution

- From the problem description we can derive the basic loop of each philosopher:

```
1. while (1) {  
2.     think();  
3.     gethashis();  
4.     eat();  
5.     puthashis();  
6. }
```

- The hashis are the shared resources that must be protected.
- Our solution must also satisfy the following requirements:
 - No deadlock
 - No starvation
 - High concurrency

Designing a solution

- To make it easier to reference the hashis, let us create two helper functions which, for a given philosopher **p**, calculate the indices of the hashis to his left and to his right.

```
1. int left(int p) { return p; }  
2. int right(int p) { return (p + 1) % 5; }
```

- We will also associate a binary semaphore with each hashi

```
1. sem_t hashis[5];  
  
2. int main(int argc, char *argv[]) {  
3.     // ...  
4.     for (int i = 0; i < 5; i++)  
5.         sem_init(&hashis[i], 0, 1);  
6.     // ...  
7. }
```


Designing `gethashis()` and `puthashis()` (version 1)

- This enables us to write a first version of the functions that implement the model

```
1. void gethashis(int p) {  
2.     sem_wait(hashis[left(p)]);  
3.     sem_wait(hashis[right(p)]);  
4. }  
  
5. void puthashis(int p) {  
6.     sem_post(hashis[left(p)]);  
7.     sem_post(hashis[right(p)]);  
8. }
```

- How do you like it? Is it correct? Does it satisfy the additional requirements?

Breaking the dependency in `gethashis()`

- We can break the circular wait in `gethashis()` by imposing a different policy for one of the philosophers (say, philosopher number 4)

```
1. void gethashis(int p) {  
2.     if (p == 4) {  
3.         sem_wait(hashis[right(p)]);  
4.         sem_wait(hashis[left(p)]);  
5.     } else {  
6.         sem_wait(hashis[left(p)]);  
7.         sem_wait(hashis[right(p)]);  
8.     }  
9. }
```

- Is that enough? Is it correct? Does it satisfy the additional requirements?