

T23a

Common Concurrency Problems


Referência principal

Ch.32 of *Operating Systems: Three Easy Pieces* by Remzi and Andrea Arpaci-Dusseau (pages.cs.wisc.edu/~remzi/OSTEP/)

Discutido em classe em 17 de outubro de 2018

Common Concurrency Problems

- More recent work focuses on studying other types of common concurrency bugs.
- Non-deadlock bugs make up 70% of concurrency bugs:
 - Atomicity violation
 - Order violation
- Deadlocks account for 30% of all bugs.



Atomicity-Violation Bugs

- In this sort of error, critical sections may not be executed atomically.
- This is a simple example found in MySQL:

```
1. Thread1::  
2.     if(thd->proc_info){  
3.         ...  
4.         fputs(thd->proc_info, ...);  
5.         ...  
6.     }  
7.  
8. Thread2::  
9.     thd->proc_info = NULL;
```

- What is the problem here?
 - Thread1 and Thread2 may access field `proc_info` in struct `thd` non-atomically.

Atomicity-Violation Bugs

- How can we solve that problem?
 - In this case it is rather straightforward: just protect the critical sections with locks.

```
1.  mutex_t lock = PTHREAD_MUTEX_INITIALIZER;

2.  Thread1::
3.      mutex_lock(&lock);
4.      if(thd->proc_info) {
5.          ...
6.          fputs(thd->proc_info, ...);
7.          ...
8.      }
9.      mutex_unlock(&lock);

10. Thread2::
11.     mutex_lock(&lock);
12.     thd->proc_info = NULL;
13.     mutex_unlock(&lock);
```

Order-Violation Bugs

- In this sort of error, instructions in different threads may be executed out of the expected order.

```
1. Thread1::  
2.     void init(){  
3.         mThread = PR_CreateThread(mMain, ...);  
4.     }  
5.  
6. Thread2::  
7.     void mMain(){  
8.         mState = mThread->State  
9.     }
```

- What is the problem here?
 - The code in Thread2 seems to assume that the variable `mThread` has already been initialized (and is not `NULL`), but this may not be true.

Order-Violation Bugs

- How can we solve that problem?
 - In this case, we can use a condition variable to enforce order of execution.

```
1. mutex_t mtLock = PTHREAD_MUTEX_INITIALIZER;
2. cond_t mtCond = PTHREAD_COND_INITIALIZER;
3. int mtInit = 0;
4.
5. Thread 1::
6. void init(){
7.     ...
8.     mThread = PR_CreateThread(mMain,...);
9.     // signal that the thread has been created.
10.    mutex_lock(&mtLock);
11.    mtInit = 1;
12.    cond_signal(&mtCond);
13.    mutex_unlock(&mtLock);
14.    ...
15. }
16. Thread2::
17. void mMain(...){
18.     ...
19.     // wait for the thread to be initialized ...
20.    mutex_lock(&mtLock);
21.    while(mtInit == 0)
22.        cond_wait(&mtCond, &mtLock);
23.    mutex_unlock(&mtLock);
24.
25.    mState = mThread->State;
26.    ...
27. }
```

Deadlocks

Deadlocks


Deadlocks...

- Are permanent blockings of two or more processes or threads that either
 - compete for system resources or
 - communicate with each other.
- Involve conflicting needs for resources by two or more processes or threads.
- Have no efficient solution.

Deadlock Definition

- Resource
 - Any (passive) thing needed by a thread to do its job
 - CPU, disk space, memory, lock
 - Preemptable
 - can be taken away by OS
 - Non-preemptable
 - must stay with thread
- Starvation
 - Thread waits indefinitely
- Deadlock
 - Circular waiting for resources
 - Deadlock \Rightarrow starvation, but not vice-versa


Bridge Crossing




- Bridge with a single lane of traffic.
- Each section of the bridge can be viewed as a resource.
- If a deadlock occurs, it can be resolved if one car backs up (preempt resources and rollback).
- Several cars may have to be backed up if a deadlock occurs.
- Starvation is possible.

Deadlock Example

Crossroads



Deadlock is possible



Deadlock has occurred

Two locks

Thread A	Thread B
<code>lock1.acquire();</code>	<code>lock2.acquire();</code>
<code>lock2.acquire();</code>	<code>lock1.acquire();</code>
<code>lock2.release();</code>	<code>lock1.release();</code>
<code>lock1.release();</code>	<code>lock2.release();</code>

- Can deadlock happen?
- Will deadlock happen?

Two locks and a condition variable


Thread A	Thread B
<code>lock1.acquire();</code>	<code>lock1.acquire();</code>
<code>...</code>	<code>...</code>
<code>lock2.acquire();</code>	<code>lock2.acquire();</code>
<code>while (needToWait()) {</code>	<code>...</code>
<code> condition.wait(lock2);</code>	<code>condition.signal(lock2);</code>
<code>}</code>	<code>...</code>
<code>lock2.release();</code>	<code>lock2.release();</code>
<code>...</code>	<code>...</code>
<code>lock1.release();</code>	<code>lock1.release();</code>

- What happens here?

Deadlock example

Dining Philosophers

- Problem
 - 5 people sit around a table and alternate between talking and eating from an infinite sushi tray placed at the center of the table.
 - There are 5 plates and 5 chopsticks as shown in the picture.
 - To eat, both chopsticks next to a plate are required.
 - Chopsticks are returned to their positions whenever a person stops or is not able to start eating.
- Can deadlock happen?
 - If so, does it always happen?
- Can a person starve?



Deadlock can arise only if 4 conditions hold simultaneously

- **Mutually exclusive finite resources**

- Resources are finite and can be used by only one process at a time.

- **Hold and wait**

- A process is holding at least one resource and is waiting to acquire additional resources held by other processes.

- **No preemption**

- A resource can be released only voluntarily by the process that is holding it.

- **Circular wait**

- There exists a set $\{P_0, P_1, \dots, P_n\}$ of waiting processes such that
 - P_0 waits for a resource that is held by P_1
 - P_1 waits for a resource that is held by P_2 ... and
 - P_n waits for a resource that is held by P_0 .

A question about the Dining Philosophers

- How does the Dining Philosophers problem meet the necessary conditions for deadlock?
 - Mutually exclusive finite resources
 - Hold and wait
 - No preemption
 - Circular wait
- Can Dining Philosophers be modified to prevent deadlock?

Methods for Handling Deadlocks

Methods for Handling Deadlocks

- **Prevention**

- Restrain requests to ensure that at least one necessary condition cannot hold so that the system can never enter a deadlock state.

- **Avoidance**

- Delay granting requests that would take the system to a state where a deadlock might occur.

- **Detection**

- Allow the system to enter a deadlock state and then recover.

- **Ignore the problem** and pretend that deadlocks never occur

- Known as the “ostrich algorithm” and used by most operating systems, including UNIX and Windows.

Deadlock Prevention:

ensure that at least one necessary condition cannot hold

■ Mutual Exclusion

- Not required for sharable resources.
- Must hold for non-sharable resources.

■ Hold and Wait

- Prevent a process from requesting a resource if it is already holding any other resources.
- Require a process to request and receive all the resources it needs before it begins execution or allow process to request resources only when it has none.
- Low resource utilization; starvation possible.

Deadlock Prevention:

ensure that at least one necessary condition cannot hold

■ No Preemption

- If a process that is holding some resources requests another resource that cannot be immediately allocated to it, then all resources currently being held are released.
- Preempted resources are added to the list of resources for which the process is waiting.
- Process will be restarted only when it can regain its old resources, as well as the new ones that it is requesting.

■ Circular Wait

- Impose a total ordering of all resource types and require that each process requests resources in an increasing order of enumeration.

Deadlock Avoidance

- In order to avoid a deadlock, the system must have some additional *a priori* information available.
 - Simplest and most useful model requires that each process declare the *maximum number* of resources of each type that it may need.
 - The deadlock-avoidance algorithm dynamically examines the resource-allocation state to ensure that there can never be a circular-wait condition.
 - Resource-allocation *state* is defined by the number of available and allocated resources, and the maximum demands of the processes.

System Model

- There are m resource types R_1, R_2, \dots, R_m
CPU cycles, memory space, I/O devices
- There are W_i instances of each resource type R_i .
- Each process uses resources in the following sequence:
 1. Request
 2. Use
 3. Release
- Requesting and releasing system resources are implemented by *system calls*.


Analysis of potential for deadlock (2 processes, 2 resources)

Process P


```
... get A;  
... get B;  
... rel A;  
... rel B;  
...
```

Process Q

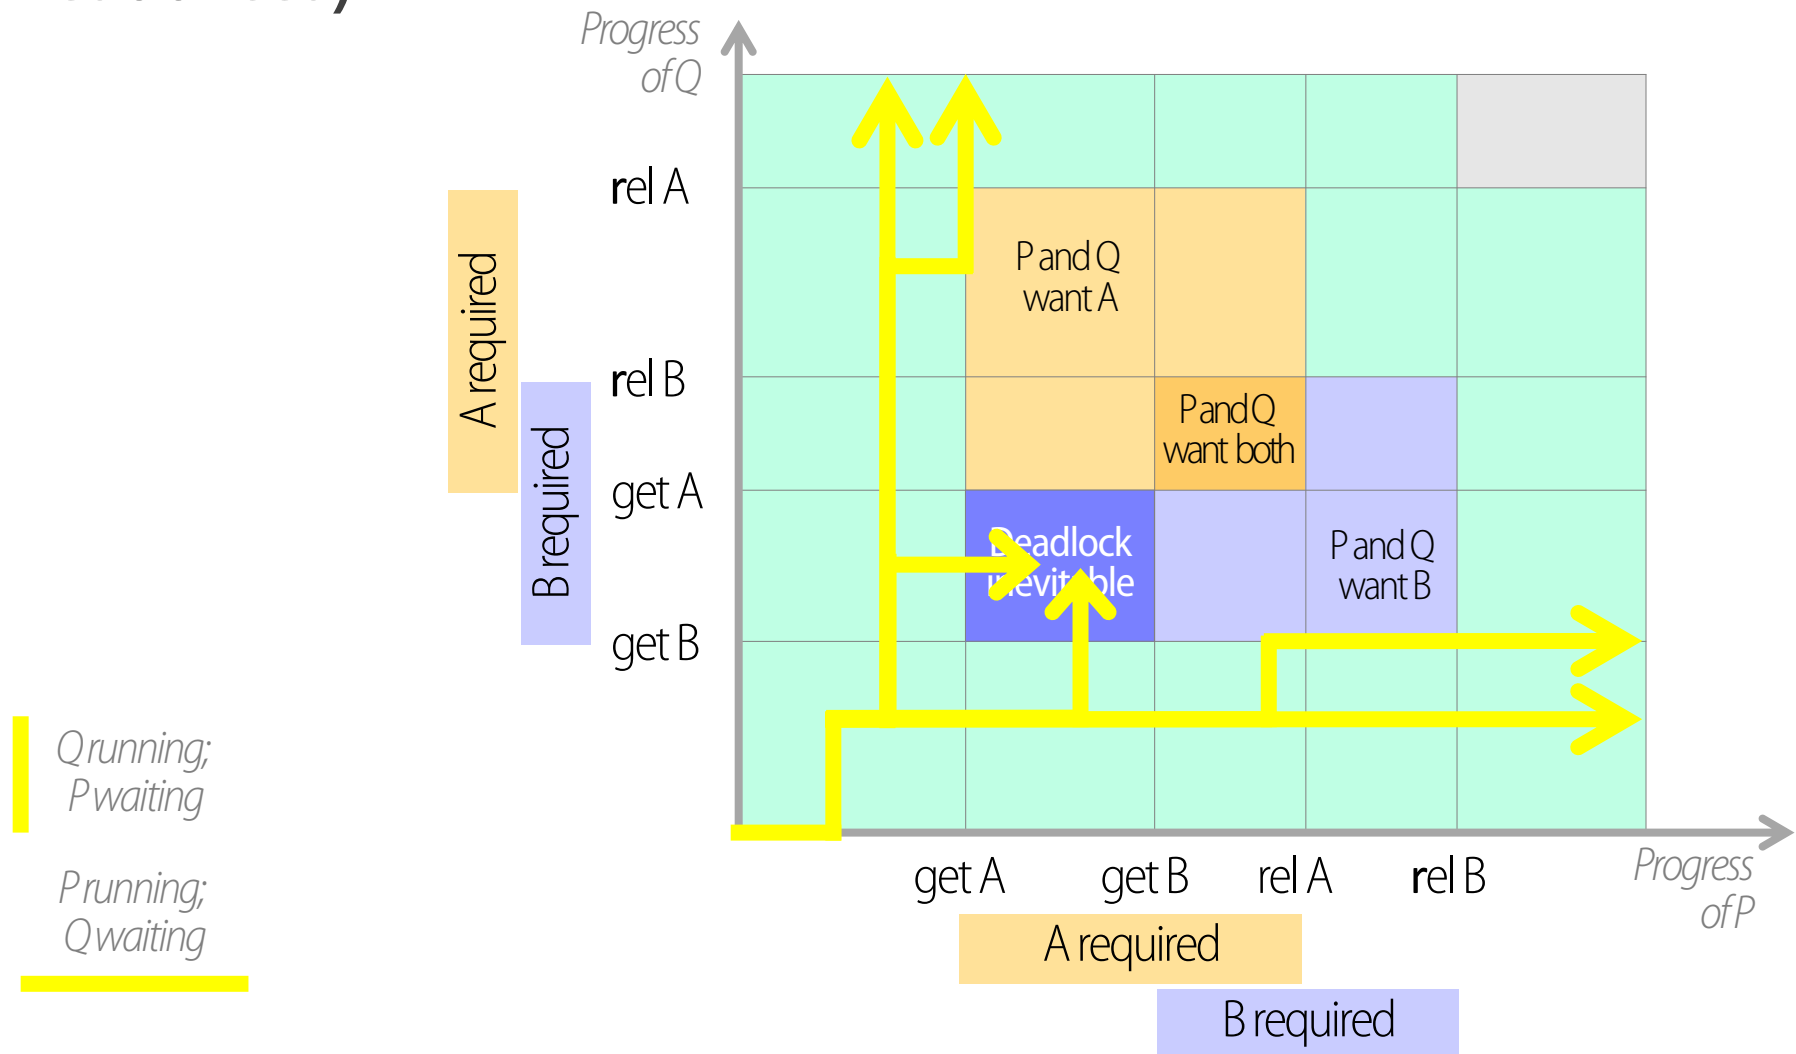
```
... get B;  
... get A;  
... rel B;  
... rel A;  
...
```



Analysis of potential for deadlock (2 processes, 2 resources)



Analysis of potential for deadlock (2 processes, 2 resources)



Deadlock avoidance: the **Safe State** concept


- When a process requests an available resource, the resource manager must decide if immediate allocation leaves the system in a safe state.
- System is said to be in a **safe state** if there exists a sequence of execution $\langle P_{i_1}, P_{i_2}, \dots, P_{i_n} \rangle$ of **all** the processes in the system such that, for each P_{i_j} , the resources that P_{i_j} can still request can be satisfied by currently available resources plus resources held by some P_{i_k} , with $k < j$.

Safe State

- In other words...
 - If P_{i_j} needs a resource that is not immediately available, then P_{i_j} can wait until all $P_{i_k}, k < j$, have finished.
 - When $P_{i_k}, k < j$, has finished, P_{i_j} can obtain the resources that it needs, execute, return allocated resources and terminate.
 - When P_{i_j} terminates, $P_{i_{j+1}}$ can obtain the resources that it needs, and so on.

Basic Facts about Safe and Unsafe States

- System is in a safe state
 - No deadlocks.
- System is in an unsafe state
 - Possibility of deadlock.
- To avoid deadlocks
 - Ensure that the system will never enter an unsafe state.



System states

Deadlock Dynamics

- Safe state
 - For any possible sequence of future resource requests, it is possible to eventually grant all requests
 - May require waiting even when resources are available!
- Unsafe state
 - Some sequence of resource requests can result in deadlock
- Doomed state
 - All possible computations lead to deadlock

Food for thought...

- What are the doomed states for Dining Philosophers?
- What are the unsafe states?
- What are the safe states?