

T23b

# Common Concurrency Problems

*Referência principal*

Ch.32 of *Operating Systems: Three Easy Pieces* by Remzi and Andrea Arpaci-Dusseau ([pages.cs.wisc.edu/~remzi/OSTEP/](http://pages.cs.wisc.edu/~remzi/OSTEP/))

*Discutido em classe em 22 de outubro de 2018*

# Deadlock avoidance: the **Safe State** concept

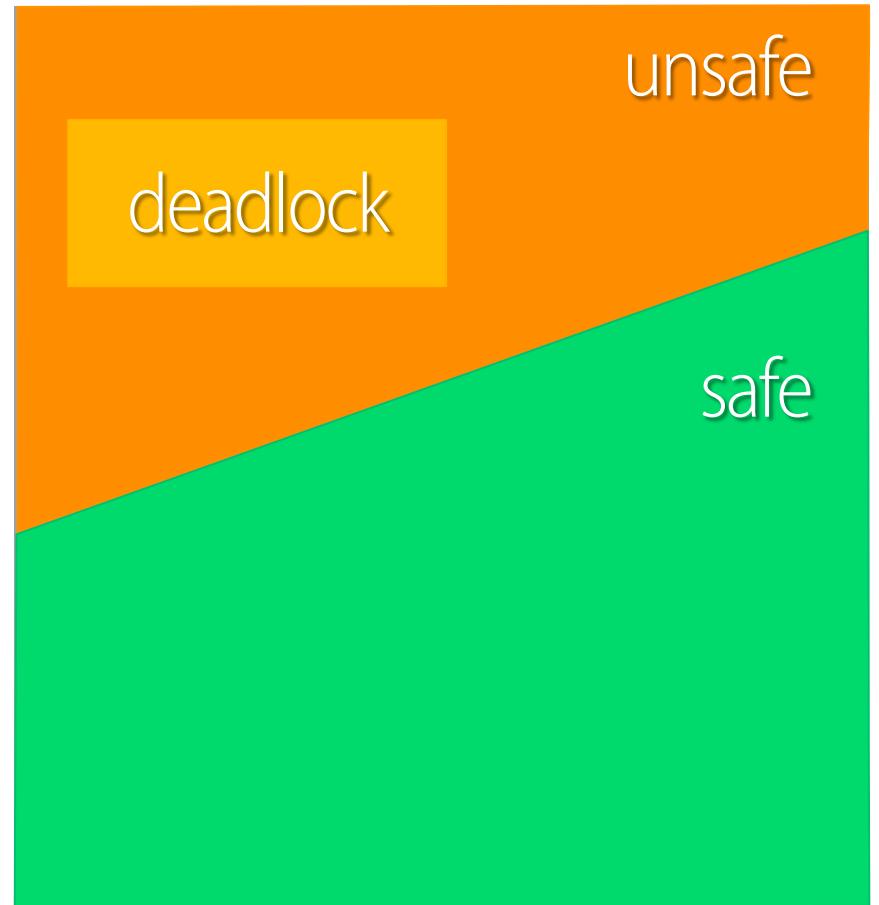
- When a process requests an available resource, the resource manager must decide if immediate allocation leaves the system in a safe state.
- A system with  $n$  processes,  $P_1, \dots, P_n$ , is said to be in a **safe state** if there exists a sequence of execution  $\langle P_{i_1}, P_{i_2}, \dots, P_{i_n} \rangle$  of all the processes in the system such that for each  $P_{i_j}$ , the resources that  $P_{i_j}$  may still need can be satisfied by currently available resources plus resources held by some  $P_{i_k}$ , with  $k < j$ .

# Safe State

- In other words...
  - If  $P_{i_j}$  needs a resource that is not immediately available, then  $P_{i_j}$  can wait until all  $P_{i_k}, k < j$ , have finished.
  - When  $P_{i_k}, k < j$ , has finished,  $P_{i_j}$  can obtain the resources that it needs, execute, return its allocated resources and terminate.
  - When  $P_{i_j}$  terminates,  $P_{i_{j+1}}$  can obtain the resources that it needs, and so on.

# Basic Facts about Safe and Unsafe States

- System is in a safe state
  - No deadlocks.
- System is in an unsafe state
  - Possibility of deadlock.
- To avoid deadlocks
  - Ensure that the system will never enter an unsafe state.



System states

# Deadlock Dynamics

- Safe state
  - For any possible sequence of future resource requests, it is possible to eventually grant all requests.
  - May require waiting even when resources are available!
- Unsafe state
  - Some sequence of resource requests can result in deadlock.
- Doomed state
  - All possible computations lead to deadlock.

# Food for thought...

- What are the doomed states for Dining Philosophers?
- What are the unsafe states?
- What are the safe states?

# Deadlock Avoidance algorithms

Single instance of a resource type

Use a resource-allocation graph scheme.

Multiple instances of a resource type

Use the banker's algorithm.

Deadlock avoidance with single instance of each resource type

# Resource-Allocation Graph Scheme

- Create a resource allocation graph.



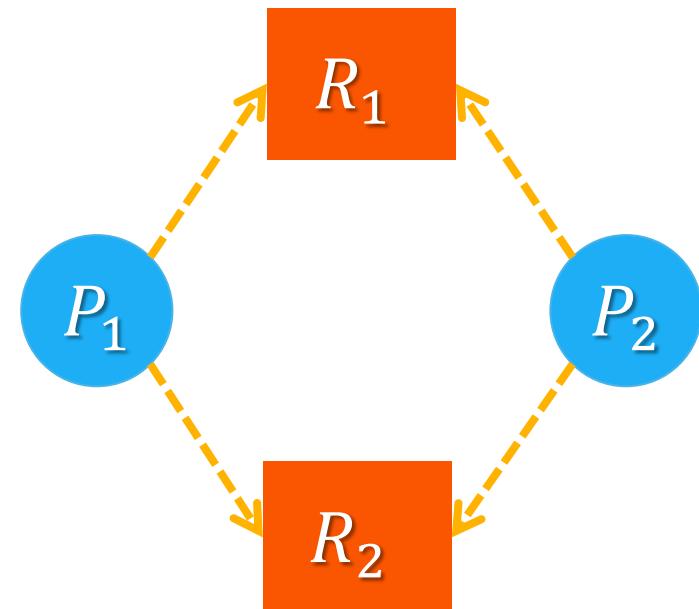
Deadlock avoidance with single instance of each resource type

# Resource-Allocation Graph Scheme

- Introduce a **claim edge**  $P_i \rightarrow R_j$

to indicate that process  $P_i$  may request resource  $R_j$  sometime.

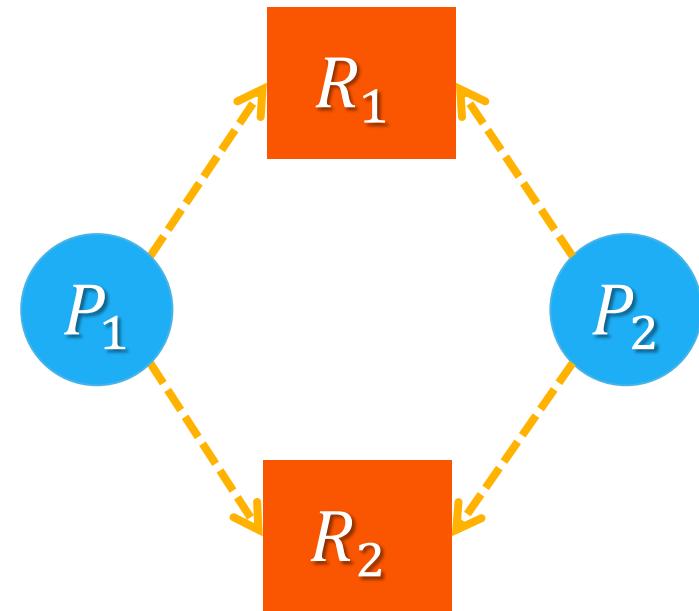
- A claim edge is denoted by a dashed line.
- Resources must be claimed *a priori*.



Deadlock avoidance with single instance of each resource type

# Resource-Allocation Graph Scheme

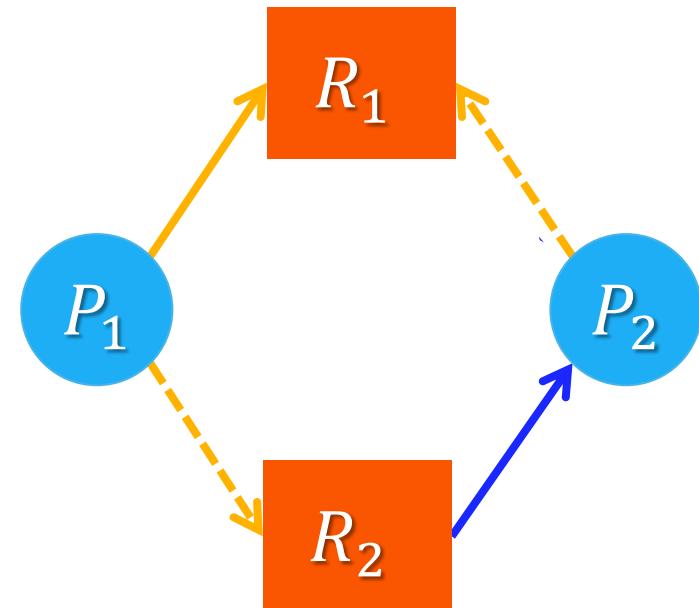
- A claim edge is converted to a **request edge** when a process requests a resource.
- A request edge is converted to an **assignment edge** when the resource is allocated to the process.
- When a resource is released by a process, assignment edge reconverts to a claim edge.



Deadlock avoidance with single instance of each resource type

# Resource-Allocation Graph Scheme

- Suppose that process  $P_i$  requests a resource  $R_j$ .
- The request can only be granted if converting the request edge to an assignment edge does not create a cycle in the graph, which would lead the system to an unsafe state.



Deadlock avoidance with multiple instances of each resource type

# Predict the Future: Banker's Algorithm

- Grant request if and only if result is a safe state.
- Sum of maximum resource needs of current threads can be greater than the total resources...
  - ... provided there is some way for all the threads to finish without getting into deadlock.
- Example
  - Proceed if and only if
    - **total available resources – number allocated  $\geq$  max remaining** that might be needed by this thread in order to finish.
    - It can be guaranteed that this thread can finish.

Deadlock avoidance with multiple instances of each resource type

# Banker's Algorithm

- Can deal with multiple instances of resources.
- Each process must claim maximum use a priori.
- When a process requests a resource it may have to wait.
- When a process gets all its resources it must return them in a finite amount of time.

# Data Structures

- Let  $P_i, 0 \leq i < n$ , be a process and  $R_k, 0 \leq k < m$ , a resource type.
- $\text{Available}_k$  (vector of length  $m$ )
  - Number of available instances of resource type  $R_k$ .
- $\text{Max}_{i,k}$  ( $n \times m$  matrix)
  - Maximum number of instances of  $R_k$  that  $P_i$  may require.
- $\text{Allocation}_{i,k}$  ( $n \times m$  matrix)
  - Number of instances of  $R_k$  currently allocated to  $P_i$ .
- $\text{Need}_{i,k}$  ( $n \times m$  matrix)
  - Number of instances of  $R_k$  that  $P_i$  may still require.
  - $\text{Need}_{i,k} = \text{Max}_{i,k} - \text{Allocation}_{i,k}$ .

# Safety Algorithm

- Let  $Work$  and  $Finish$  be vectors of length  $m$  and  $n$ , respectively.
  1. Initially
$$Work = Available$$
$$Finish_i = false, 0 \leq i < n$$
  2. Find an  $i$  such that  $Finish_i = false$  and  $Need_i \leq Work$ .  
If no such  $i$  exists, go to step 4.
  3. Let
$$Work \leftarrow Work + Allocation_i$$
$$Finish_i \leftarrow true$$
and go back to step 2.
  4. If  $Finish_i = true$  for all  $0 \leq i < n$ , the system is in a safe state.

# Resource-Request Algorithm for Process $P_i$

1. Let  $\text{Request}$  be the request vector for process  $P_i$ .
  - $\text{Request}_k = x$ , means that  $P_i$  wants  $x$  instances more of resource  $R_k$ .
2. If  $\text{Request} > \text{Need}_i$ , raise error condition, since  $P_i$  has exceeded its maximum claim.
3. If  $\text{Request} > \text{Available}$ ,  $P_i$  must wait, since resources are not immediately available.
4. ...

# Resource-Request Algorithm for Process $P_i$

4. Pretend to allocate the requested resources to  $P_i$ :
  - $\text{Available} \leftarrow \text{Available} - \text{Request}$
  - $\text{Allocation}_i \leftarrow \text{Allocation}_i + \text{Request}$
  - $\text{Need}_i \leftarrow \text{Need}_i - \text{Request}$
5. Run the safety algorithm.
  - If state is safe, allocate the requested resources to  $P_i$ .
  - If state is unsafe,  $P_i$  must wait and the original state of resource allocation is restored.

# Example of Banker's Algorithm

- 5 processes  $P_0$  through  $P_4$
- 3 resource types,  
A (10 instances), B (5), C (7).

Total		
A	B	C
10	5	7

System snapshot at time  $T_0$

	Max		
	A	B	C
$P_0$	7	5	3
$P_1$	3	2	2
$P_2$	9	0	2
$P_3$	2	2	2
$P_4$	4	3	3

	Allocation		
	A	B	C
$P_0$			
$P_1$			
$P_2$			
$P_3$			
$P_4$			

	Need		
	A	B	C
$P_0$			
$P_1$			
$P_2$			
$P_3$			
$P_4$			

Available		
A	B	C

# Example of Banker's Algorithm

- 5 processes  $P_0$  through  $P_4$
- 3 resource types,  
A (10 instances), B (5), C (7).

Total		
A	B	C
10	5	7

System snapshot at time  $T_0$

	Max		
	A	B	C
$P_0$	7	5	3
$P_1$	3	2	2
$P_2$	9	0	2
$P_3$	2	2	2
$P_4$	4	3	3

	Allocation		
	A	B	C
$P_0$	0	1	0
$P_1$	2	0	0
$P_2$	3	0	2
$P_3$	2	1	1
$P_4$	0	0	2

	Need		
	A	B	C
$P_0$	7	4	3
$P_1$	1	2	2
$P_2$	6	0	0
$P_3$	0	1	1
$P_4$	4	3	3

	Available		
	A	B	C
	10	5	7

# Example of Banker's Algorithm

- 5 processes  $P_0$  through  $P_4$
- 3 resource types,  
A (10 instances), B (5), C (7).

Total		
A	B	C
10	5	7

System snapshot at time  $T_0$

	Max		
	A	B	C
$P_0$	7	5	3
$P_1$	3	2	2
$P_2$	9	0	2
$P_3$	2	2	2
$P_4$	4	3	3

	Allocation		
	A	B	C
$P_0$	0	1	0
$P_1$	2	0	0
$P_2$	3	0	2
$P_3$	2	1	1
$P_4$	0	0	2

	Need		
	A	B	C
$P_0$	7	4	3
$P_1$	1	2	2
$P_2$	6	0	0
$P_3$	0	1	1
$P_4$	4	3	3

	Available		
	A	B	C
	3	3	2

# Example of Banker's Algorithm

- 5 processes  $P_0$  through  $P_4$
- 3 resource types,  
A (10 instances), B (5), C (7).

Total		
A	B	C
10	5	7

System snapshot at time  $T_0$

	Max		
	A	B	C
$P_0$	7	5	3
$P_1$	3	2	2
$P_2$	9	0	2
$P_3$	2	2	2
$P_4$	4	3	3

	Allocation		
	A	B	C
$P_0$	0	1	0
$P_1$	2	0	0
$P_2$	3	0	2
$P_3$	2	1	1
$P_4$	0	0	2

	Need		
	A	B	C
$P_0$	7	4	3
$P_1$			
$P_2$			
$P_3$			
$P_4$			

	Available		
	A	B	C
	3	3	2

# Example of Banker's Algorithm

- 5 processes  $P_0$  through  $P_4$
- 3 resource types,  
A (10 instances), B (5), C (7).

Total		
A	B	C
10	5	7

System snapshot at time  $T_0$

	Max		
	A	B	C
$P_0$	7	5	3
$P_1$	3	2	2
$P_2$	9	0	2
$P_3$	2	2	2
$P_4$	4	3	3

	Allocation		
	A	B	C
$P_0$	0	1	0
$P_1$	2	0	0
$P_2$	3	0	2
$P_3$	2	1	1
$P_4$	0	0	2

	Need		
	A	B	C
$P_0$	7	4	3
$P_1$	1	2	2
$P_2$			
$P_3$			
$P_4$			

Available		
A	B	C
3	3	2

# Example of Banker's Algorithm

- 5 processes  $P_0$  through  $P_4$
- 3 resource types,  
A (10 instances), B (5), C (7).

Total		
A	B	C
10	5	7

System snapshot at time  $T_0$

	Max		
	A	B	C
$P_0$	7	5	3
$P_1$	3	2	2
$P_2$	9	0	2
$P_3$	2	2	2
$P_4$	4	3	3

	Allocation		
	A	B	C
$P_0$	0	1	0
$P_1$	2	0	0
$P_2$	3	0	2
$P_3$	2	1	1
$P_4$	0	0	2

	Need		
	A	B	C
$P_0$	7	4	3
$P_1$	1	2	2
$P_2$	6	0	0
$P_3$	0	1	1
$P_4$	4	3	1

	Available		
	A	B	C
	3	3	2

# Example of Banker's Algorithm

- Is the system state at time  $T_0$  safe?
  - To find the answer we must run the safety algorithm.

	Max			Allocation			Need			Available		
	A	B	C	A	B	C	A	B	C	A	B	C
P <sub>0</sub>	7	5	3	0	1	0	7	4	3	3	3	2
P <sub>1</sub>	3	2	2	2	0	0	1	2	2			
P <sub>2</sub>	9	0	2	3	0	2	6	0	0			
P <sub>3</sub>	2	2	2	2	1	1	0	1	1			
P <sub>4</sub>	4	3	3	0	0	2	4	3	1			

	Work		
	A	B	C

# Example of Banker's Algorithm

- Is the system state at time  $T_0$  safe?
  - First, create *Work* and *Finish* vectors.

	Max			Allocation			Need			Finish			Available		
	A	B	C	A	B	C	A	B	C				A	B	C
P <sub>0</sub>	7	5	3	0	1	0	7	4	3				3	3	2
P <sub>1</sub>	3	2	2	2	0	0	1	2	2						
P <sub>2</sub>	9	0	2	3	0	2	6	0	0						
P <sub>3</sub>	2	2	2	2	1	1	0	1	1						
P <sub>4</sub>	4	3	3	0	0	2	4	3	1						

	Work		
	A	B	C
	3	3	2

# Example of Banker's Algorithm

- Is the system state at time  $T_0$  safe?
    - Can any process terminate?
      - Yes.  $P_1$  can. Retrieve the resources that it holds.

	Max		
	A	B	C
P <sub>0</sub>	7	5	3
P <sub>1</sub>	3	2	2
P <sub>2</sub>	9	0	2
P <sub>3</sub>	2	2	2
P <sub>4</sub>	4	3	3

Allocation		
A	B	C
0	1	0
2	0	0
3	0	2
2	1	1
0	0	2

Need		
A	B	C
7	4	3
1	2	2
6	0	0
0	1	1
4	3	1

# Finish

Available		
A	B	C
3	3	2

Work		
A	B	C
5	3	2

# Example of Banker's Algorithm

- Is the system state at time  $T_0$  safe?
    - Can another process terminate?
      - Yes.  $P_3$  can. Retrieve the resources that it holds.

	Max		
	A	B	C
P <sub>0</sub>	7	5	3
P <sub>1</sub>	3	2	2
P <sub>2</sub>	9	0	2
P <sub>3</sub>	2	2	2
P <sub>4</sub>	4	3	3

Allocation		
A	B	C
0	1	0
2	0	0
3	0	2
2	1	1
0	0	2

Need		
A	B	C
7	4	3
1	2	2
6	0	0
0	1	1
4	3	1

Finish  
T  
T

Available		
A	B	C
3	3	2
Work		
A	B	C
7	4	3

# Example of Banker's Algorithm

- Is the system state at time  $T_0$  safe?
    - Now,  $P_0$  can terminate. Let's retrieve its resources.

	Max				
	A	B	C		
P <sub>0</sub>	7	5	3		
P <sub>1</sub>	3	2	2		
P <sub>2</sub>	9	0	2		
P <sub>3</sub>	2	2	2		
P <sub>4</sub>	4	3	3		

	Allocation				
	A	B	C		
P <sub>0</sub>	0	1	0		
P <sub>1</sub>	2	0	0		
P <sub>2</sub>	3	0	2		
P <sub>3</sub>	2	1	1		
P <sub>4</sub>	0	0	2		

	Need				
	A	B	C		
P <sub>0</sub>	7	4	3		
P <sub>1</sub>	1	2	2		
P <sub>2</sub>	6	0	0		
P <sub>3</sub>	0	1	1		
P <sub>4</sub>	4	3	1		

	Finish				
	A	B	C		
P <sub>0</sub>	T				
P <sub>1</sub>	T				
P <sub>2</sub>					
P <sub>3</sub>					
P <sub>4</sub>					

	Available				
	A	B	C		
P <sub>0</sub>	3	3	2		
P <sub>1</sub>					
P <sub>2</sub>					
P <sub>3</sub>					
P <sub>4</sub>					

	Work				
	A	B	C		
P <sub>0</sub>	7	5	3		
P <sub>1</sub>					
P <sub>2</sub>					
P <sub>3</sub>					
P <sub>4</sub>					

# Example of Banker's Algorithm

- Is the system state at time  $T_0$  safe?
  - Now,  $P_2$  can terminate.

	Max			Allocation	Need	Finish	Available	Work		
	A	B	C	A	B	C	A	A	B	C
$P_0$	7	5	3	0	1	0	7	4	3	T
$P_1$	3	2	2	2	0	0	1	2	2	T
$P_2$	9	0	2	3	0	2	6	0	0	T
$P_3$	2	2	2	2	1	1	0	1	1	T
$P_4$	4	3	3	0	0	2	4	3	1	

# Example of Banker's Algorithm

- Is the system state at time  $T_0$  safe?
  - And finally  $P_4$  can terminate.
  - YES!

	Max			Allocation			Need			Finish			Available		
	A	B	C	A	B	C	A	B	C				A	B	C
$P_0$	7	5	3	0	1	0	7	4	3	T			3	3	2
$P_1$	3	2	2	2	0	0	1	2	2	T					
$P_2$	9	0	2	3	0	2	6	0	0	T					
$P_3$	2	2	2	2	1	1	0	1	1	T					
$P_4$	4	3	3	0	0	2	4	3	1	T					
	Work														
	A	B	C							A	B	C	10	5	7

# Example of Banker's Algorithm

- Can a (1, 0, 2) request by  $P_1$  be granted at  $T_0$ ?
  - Pretend to grant the request

	Max						
	A	B	C				
$P_0$	7	5	3				
$P_1$	3	2	2	0	1	0	
$P_2$	9	0	2	2	0	0	
$P_3$	2	2	2	3	0	2	
$P_4$	4	3	3	2	1	1	

	Allocation						
	A	B	C				
$P_0$	0	1	0				
$P_1$	2	0	0	2	0	0	
$P_2$	3	0	2	3	0	2	
$P_3$	2	1	1	2	1	1	
$P_4$	0	0	2	0	0	2	

	Need						
	A	B	C				
$P_0$	7	4	3	7	4	3	
$P_1$	1	2	2	1	2	2	
$P_2$	6	0	0	6	0	0	
$P_3$	0	1	1	0	1	1	
$P_4$	4	3	1	4	3	1	

	Finish						
	A	B	C				
$P_0$							
$P_1$							
$P_2$							
$P_3$							
$P_4$							

	Available						
	A	B	C				
$P_0$	3	3	2	3	3	2	
$P_1$							
$P_2$							
$P_3$							
$P_4$							

	Work						
	A	B	C				
$P_0$							
$P_1$							
$P_2$							
$P_3$							
$P_4$							

# Example of Banker's Algorithm

- At  $T_0$  can a  $(1, 0, 2)$  request by  $P_1$  be granted?
  - Pretend to grant the request
  - Run the safety algorithm

	Max		
	A	B	C
$P_0$	7	5	3
$P_1$	3	2	2
$P_2$	9	0	2
$P_3$	2	2	2
$P_4$	4	3	3

	Allocation		
	A	B	C
	0	1	0
$P_1$	3	0	2
	3	0	2
	2	1	1
	0	0	2

	Need		
	A	B	C
	7	4	3
$P_1$	0	2	0
	6	0	0
	0	1	1
	4	3	1

	Finish		

	Available		
	A	B	C
	2	3	0

	Work		
	A	B	C

# Example of Banker's Algorithm

- At  $T_0$  can a  $(1, 0, 2)$  request by  $P_1$  be granted?
  - Pretend to grant the request
  - Run the safety algorithm

	Max		
	A	B	C
$P_0$	7	5	3
$P_1$	3	2	2
$P_2$	9	0	2
$P_3$	2	2	2
$P_4$	4	3	3

	Allocation		
	A	B	C
	0	1	0
$P_1$	3	0	2
$P_2$	3	0	2
$P_3$	2	1	1
$P_4$	0	0	2

	Need		
	A	B	C
	7	4	3
$P_1$	0	2	0
$P_2$	6	0	0
$P_3$	0	1	1
$P_4$	4	3	1

	Finish		

	Available		
	A	B	C
	2	3	0

	Work		
	A	B	C
	2	3	0

# Example of Banker's Algorithm

- At  $T_0$  can a  $(1, 0, 2)$  request by  $P_1$  be granted?
  - Pretend to grant the request
  - Run the safety algorithm

	Max		
	A	B	C
$P_0$	7	5	3
$P_1$	3	2	2
$P_2$	9	0	2
$P_3$	2	2	2
$P_4$	4	3	3

	Allocation		
	A	B	C
	0	1	0
$P_1$	3	0	2
$P_2$	3	0	2
$P_3$	2	1	1
$P_4$	0	0	2

	Need		
	A	B	C
	7	4	3
$P_1$	0	2	0
$P_2$	6	0	0
$P_3$	0	1	1
$P_4$	4	3	1

	Finish		
$T$			

	Available		
	A	B	C
	2	3	0

	Work		
	A	B	C
	5	3	2

# Example of Banker's Algorithm

- At  $T_0$  can a  $(1, 0, 2)$  request by  $P_1$  be granted?
  - Pretend to grant the request
  - Run the safety algorithm

	Max			Allocation			Need			Available		
	A	B	C	A	B	C	A	B	C	A	B	C
$P_0$	7	5	3	0	1	0	7	4	3	2	3	0
$P_1$	3	2	2	3	0	2	0	2	0	T		
$P_2$	9	0	2	3	0	2	6	0	0			
$P_3$	2	2	2	2	1	1	0	1	1	T		
$P_4$	4	3	3	0	0	2	4	3	1			

	Work		
	A	B	C
	7	4	3

# Example of Banker's Algorithm

- At  $T_0$  can a  $(1, 0, 2)$  request by  $P_1$  be granted?
  - Pretend to grant the request
  - Run the safety algorithm

	Max		
	A	B	C
$P_0$	7	5	3
$P_1$	3	2	2
$P_2$	9	0	2
$P_3$	2	2	2
$P_4$	4	3	3

	Allocation		
	A	B	C
	0	1	0
$P_1$	3	0	2
$P_2$	3	0	2
$P_3$	2	1	1
$P_4$	0	0	2

	Need		
	A	B	C
	7	4	3
$P_1$	0	2	0
$P_2$	6	0	0
$P_3$	0	1	1
$P_4$	4	3	1

	Finish		
	A	B	C
	T		
$P_1$	T		
$P_2$			
$P_3$	T		
$P_4$			

	Available		
	A	B	C
	2	3	0

	Work		
	A	B	C
	7	5	3

# Example of Banker's Algorithm

- At  $T_0$  can a  $(1, 0, 2)$  request by  $P_1$  be granted?
  - Pretend to grant the request
  - Run the safety algorithm

	Max		
	A	B	C
$P_0$	7	5	3
$P_1$	3	2	2
$P_2$	9	0	2
$P_3$	2	2	2
$P_4$	4	3	3

	Allocation		
	A	B	C
	0	1	0
	3	0	2
	3	0	2
	2	1	1
	0	0	2

	Need		
	A	B	C
	7	4	3
	0	2	0
	6	0	0
	0	1	1
	4	3	1

	Finish		
	T		
	T		
	T		
	T		

	Available		
	A	B	C
	2	3	0

	Work		
	A	B	C
	10	5	5

# Example of Banker's Algorithm

- At  $T_0$  can a  $(1, 0, 2)$  request by  $P_1$  be granted?
    - Pretend to grant the request
    - Run the safety algorithm

	Max		
	A	B	C
P <sub>0</sub>	7	5	3
P <sub>1</sub>	3	2	2
P <sub>2</sub>	9	0	2
P <sub>3</sub>	2	2	2
P <sub>4</sub>	4	3	3

Allocation		
A	B	C
0	1	0
3	0	2
3	0	2
2	1	1
0	0	2

Need		
A	B	C
7	4	3
0	2	0
6	0	0
0	1	1
4	3	1

Finish  
T  
T  
T  
T  
T

Available		
A	B	C
2	3	0

Work		
A	B	C
10	5	7

# Example of Banker's Algorithm

- At  $T_0$  can a  $(1, 0, 2)$  request by  $P_1$  be granted?

- Pretend to grant the request
- Run the safety algorithm

→ Yes!

	Max		
	A	B	C
$P_0$	7	5	3
$P_1$	3	2	2
$P_2$	9	0	2
$P_3$	2	2	2
$P_4$	4	3	3

	Allocation		
	A	B	C
	0	1	0
$P_1$	3	0	2
$P_2$	3	0	2
$P_3$	2	1	1
$P_4$	0	0	2

	Need		
	A	B	C
	7	4	3
$P_1$	0	2	0
$P_2$	6	0	0
$P_3$	0	1	1
$P_4$	4	3	1

	Finish		
	T		
$P_1$	T		
$P_2$	T		
$P_3$	T		
$P_4$	T		

	Available		
	A	B	C
	2	3	0

	Work		
	A	B	C
	10	5	7

# Example of Banker's Algorithm

- At  $T_0$  can a  $(3, 3, 0)$  request by  $P_4$  be granted?
  - Pretend to grant the request

	Max			Allocation			Need			Finish			Available		
	A	B	C	A	B	C	A	B	C				A	B	C
$P_0$	7	5	3	0	1	0	7	4	3				3	3	2
$P_1$	3	2	2	2	0	0	1	2	2						
$P_2$	9	0	2	3	0	2	6	0	0						
$P_3$	2	2	2	2	1	1	0	1	1						
$P_4$	4	3	3	0	0	2	4	3	1						

	Work		
	A	B	C
	3	3	2

# Example of Banker's Algorithm

- At  $T_0$  can a  $(3, 3, 0)$  request by  $P_4$  be granted?
  - Pretend to grant the request

	Max			Allocation			Need			Finish			Available		
	A	B	C	A	B	C	A	B	C				A	B	C
$P_0$	7	5	3	0	1	0	7	4	3				0	0	2
$P_1$	3	2	2	2	0	0	1	2	2						
$P_2$	9	0	2	3	0	2	6	0	0						
$P_3$	2	2	2	2	1	1	0	1	1						
$P_4$	4	3	3	3	3	2	1	0	1				0	0	2
	Work														
	A	B	C										A	B	C
	0	0	2										0	0	2

# Example of Banker's Algorithm

- At  $T_0$  can a  $(3, 3, 0)$  request by  $P_4$  be granted?

- Pretend to grant the request
- Run the safety algorithm

→ No!

	Max		
	A	B	C
$P_0$	7	5	3
$P_1$	3	2	2
$P_2$	9	0	2
$P_3$	2	2	2
$P_4$	4	3	3

	Allocation		
	A	B	C
	0	1	0
	2	0	0
	3	0	2
	2	1	1
	3	3	2

	Need		
	A	B	C
	7	4	3
	1	2	2
	6	0	0
	0	1	1
	1	0	1

	Finish		

	Available		
	A	B	C
	0	0	2

	Work		
	A	B	C
	0	0	2

# Example of Banker's Algorithm

- At  $T_0$  can a  $(0, 2, 0)$  request by  $P_0$  be granted?

Max		Allocation			Need			Available		
	A	B	C	A	B	C	A	B	C	
$P_0$	7	5	3	0	1	0	7	4	3	
$P_1$	3	2	2	2	0	0	1	2	2	
$P_2$	9	0	2	3	0	2	6	0	0	
$P_3$	2	2	2	2	1	1	0	1	1	
$P_4$	4	3	3	0	0	2	4	3	1	

Work		
A	B	C

# Example of Banker's Algorithm

- At  $T_0$  can a  $(0, 2, 0)$  request by  $P_0$  be granted?
    - Pretend to grant the request

	Max		
	A	B	C
P <sub>0</sub>	7	5	3
P <sub>1</sub>	3	2	2
P <sub>2</sub>	9	0	2
P <sub>3</sub>	2	2	2
P <sub>4</sub>	4	3	3

Allocation		
A	B	C
0	1	0
2	0	0
3	0	2
2	1	1
0	0	2

Need		
A	B	C
7	4	3
1	2	2
6	0	0
0	1	1
4	3	1

Finish

Available		
A	B	C
3	3	2

Work		
A	B	C

# Example of Banker's Algorithm

- At  $T_0$  can a  $(0, 2, 0)$  request by  $P_0$  be granted?
  - Pretend to grant the request

	Max						
	A	B	C				
$P_0$	7	5	3				
$P_1$	3	2	2	2	0	0	
$P_2$	9	0	2	3	0	2	
$P_3$	2	2	2	2	1	1	
$P_4$	4	3	3	0	0	2	

	Allocation						
	A	B	C				
$P_0$	0	3	0	0	3	0	
$P_1$	2	0	0	2	0	0	
$P_2$	3	0	2	3	0	2	
$P_3$	2	1	1	2	1	1	
$P_4$	0	0	2	0	0	2	

	Need						
	A	B	C				
$P_0$	7	2	3	7	2	3	
$P_1$	1	2	2	1	2	2	
$P_2$	6	0	0	6	0	0	
$P_3$	0	1	1	0	1	1	
$P_4$	4	3	1	4	3	1	

	Finish						
	A	B	C				
$P_0$							
$P_1$							
$P_2$							
$P_3$							
$P_4$							

	Available						
	A	B	C				
$P_0$	3	1	2	3	1	2	
$P_1$							
$P_2$							
$P_3$							
$P_4$							

	Work						
	A	B	C				
$P_0$							
$P_1$							
$P_2$							
$P_3$							
$P_4$							

# Example of Banker's Algorithm

- At  $T_0$  can a  $(0, 2, 0)$  request by  $P_0$  be granted?
  - Pretend to grant the request
  - Run the safety algorithm

	Max			Allocation			Need			Available			Work		
	A	B	C	A	B	C	A	B	C	A	B	C	A	B	C
$P_0$	7	5	3	0	3	0	7	2	3	3	1	2	3	1	2
$P_1$	3	2	2	2	0	0	1	2	2						
$P_2$	9	0	2	3	0	2	6	0	0						
$P_3$	2	2	2	2	1	1	0	1	1						
$P_4$	4	3	3	0	0	2	4	3	1						

# Example of Banker's Algorithm

- At  $T_0$  can a  $(0, 2, 0)$  request by  $P_0$  be granted?
  - Pretend to grant the request
  - Run the safety algorithm

	Max			Allocation			Need			Finish			Available		
	A	B	C	A	B	C	A	B	C				A	B	C
$P_0$	7	5	3	0	3	0	7	2	3				3	1	2
$P_1$	3	2	2	2	0	0	1	2	2						
$P_2$	9	0	2	3	0	2	6	0	0						
$P_3$	2	2	2	2	1	1	0	1	1	T					
$P_4$	4	3	3	0	0	2	4	3	1				5	2	3

# Example of Banker's Algorithm

- At  $T_0$  can a  $(0, 2, 0)$  request by  $P_0$  be granted?
    - Pretend to grant the request
    - Run the safety algorithm

	Max		
	A	B	C
P <sub>0</sub>	7	5	3
P <sub>1</sub>	3	2	2
P <sub>2</sub>	9	0	2
P <sub>3</sub>	2	2	2
P <sub>4</sub>	4	3	3

Allocation		
A	B	C
0	3	0
2	0	0
3	0	2
2	1	1
0	0	2

Need		
A	B	C
7	2	3
1	2	2
6	0	0
0	1	1
4	3	1

A vertical stack of five rectangular cards. The top card is orange and labeled "Finish". The bottom four cards are yellow and each have a large black letter "T" on them.

Available		
A	B	C
3	1	2

Work		
A	B	C
7	2	3

# Example of Banker's Algorithm

- At  $T_0$  can a  $(0, 2, 0)$  request by  $P_0$  be granted?
    - Pretend to grant the request
    - Run the safety algorithm

	Max		
	A	B	C
P <sub>0</sub>	7	5	3
P <sub>1</sub>	3	2	2
P <sub>2</sub>	9	0	2
P <sub>3</sub>	2	2	2
P <sub>4</sub>	4	3	3

Allocation		
A	B	C
0	3	0
2	0	0
3	0	2
2	1	1
0	0	2

Need		
A	B	C
7	2	3
1	2	2
6	0	0
0	1	1
4	3	1

A vertical stack of five rectangular cards. The top card is orange and labeled "Finish". The bottom four cards are yellow and each have a large black letter "T" on them.

Available		
A	B	C
3	1	2

Work		
A	B	C
7	5	3

# Example of Banker's Algorithm

- At  $T_0$  can a  $(0, 2, 0)$  request by  $P_0$  be granted?
    - Pretend to grant the request
    - Run the safety algorithm

	Max		
	A	B	C
P <sub>0</sub>	7	5	3
P <sub>1</sub>	3	2	2
P <sub>2</sub>	9	0	2
P <sub>3</sub>	2	2	2
P <sub>4</sub>	4	3	3

Allocation		
A	B	C
0	3	0
2	0	0
3	0	2
2	1	1
0	0	2

Need		
A	B	C
7	2	3
1	2	2
6	0	0
0	1	1
4	3	1

Finish  
T  
T  
T  
T

Available		
A	B	C
3	1	2

Work		
A	B	C
10	5	5

# Example of Banker's Algorithm

- At  $T_0$  can a  $(0, 2, 0)$  request by  $P_0$  be granted?

- Pretend to grant the request
- Run the safety algorithm

→ Yes!

	Max		
	A	B	C
$P_0$	7	5	3
$P_1$	3	2	2
$P_2$	9	0	2
$P_3$	2	2	2
$P_4$	4	3	3

	Allocation		
	A	B	C
$P_0$	0	3	0
$P_1$	2	0	0
$P_2$	3	0	2
$P_3$	2	1	1
$P_4$	0	0	2

	Need		
	A	B	C
$P_0$	7	2	3
$P_1$	1	2	2
$P_2$	6	0	0
$P_3$	0	1	1
$P_4$	4	3	1

	Finish		
$P_0$	T		
$P_1$		T	
$P_2$			T
$P_3$			T
$P_4$			T

	Available		
	A	B	C
	3	1	2

	Work		
	A	B	C
	10	5	7

# Deadlock Detection

- Allow system to enter deadlock state
- Run a detection algorithm
  - Resource allocation graphs are used to detect the existence or possibility of deadlock.
- Apply a recovery scheme

# Detect and Repair

- Algorithm
  - Scan wait-for graph
  - Detect cycles
  - Fix cycles
- Proceed without the resource
  - Requires robust exception handling code
- Roll back and retry
  - Transaction: all operations are provisional until have all required resources to complete operation

# Resource-Allocation Graph (RAG)

- A RAG is a directed graph given by a set of vertices  $V$  and a set of edges  $E$  such that
  - $V$  is partitioned into two subsets:
    - $P = \{P_1, P_2, \dots, P_n\}$ , the set consisting of all the processes in the system
    - $R = \{R_1, R_2, \dots, R_m\}$ , the set consisting of all resource types in the system.
  - A directed edge  $P_i \rightarrow R_k$  is called a *request edge*.
  - A directed edge  $R_k \rightarrow P_i$  is called an *assignment edge*.

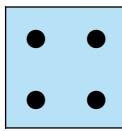
## Deadlock Detection

# Resource-Allocation Graphs

- Process  $P_i$

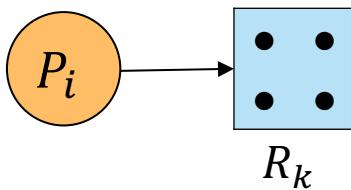


- Resource Type  $R_k$  with 4 instances

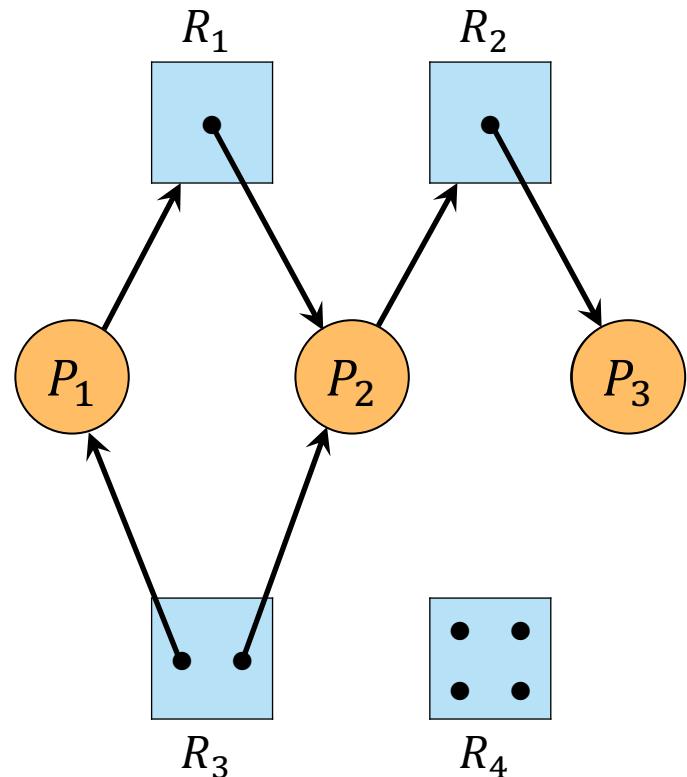
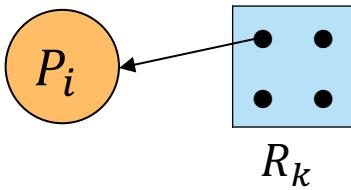


$R_k$

- $P_i$  requests an instance of  $R_k$

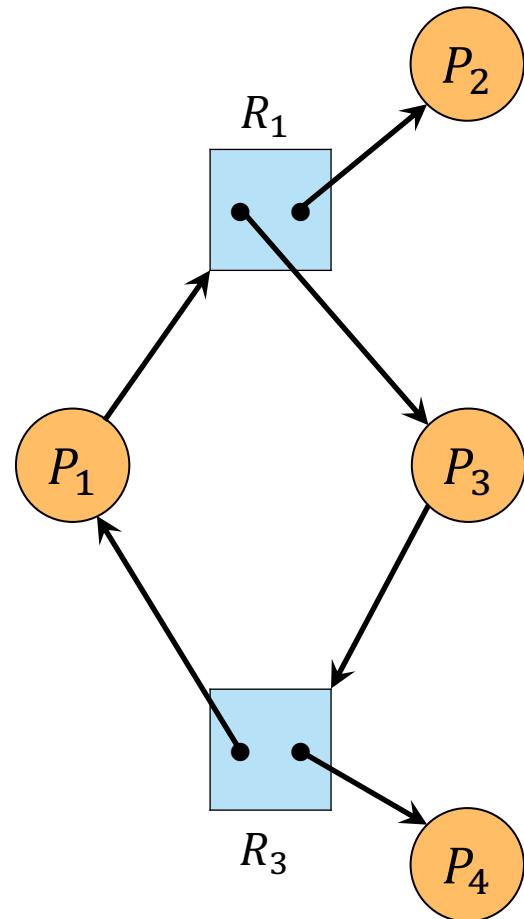
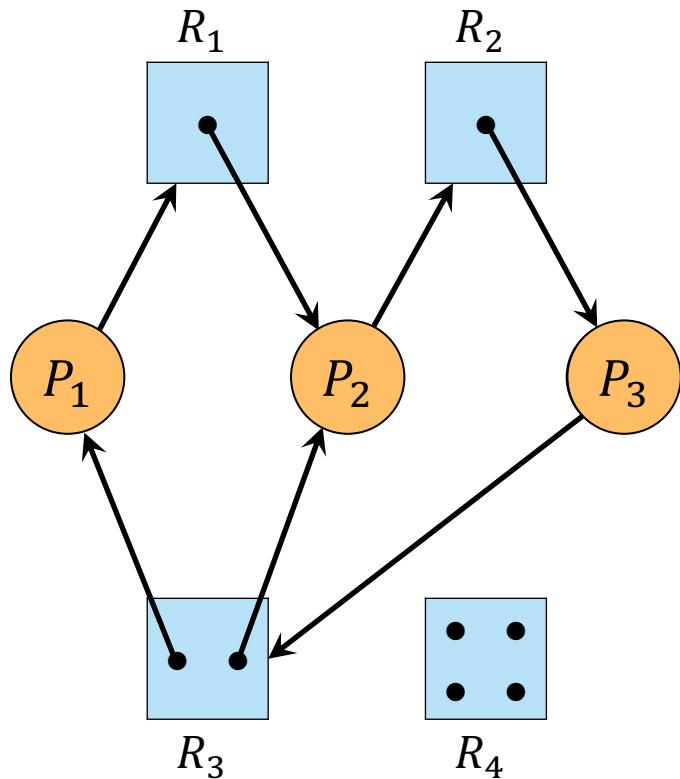


- $P_i$  holds an instance of  $R_k$



Deadlock Detection

# Do these Resource Allocation Graphs show deadlocks?

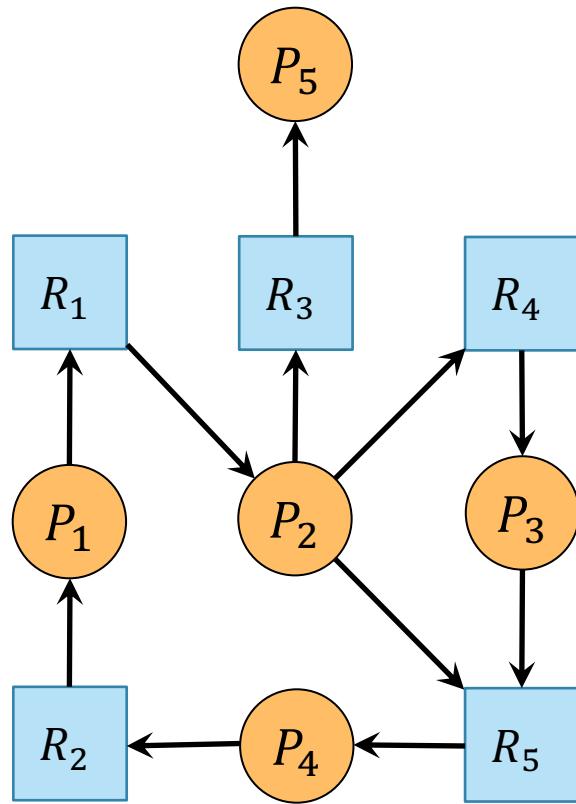


# Basic facts about a Resource Allocation Graph

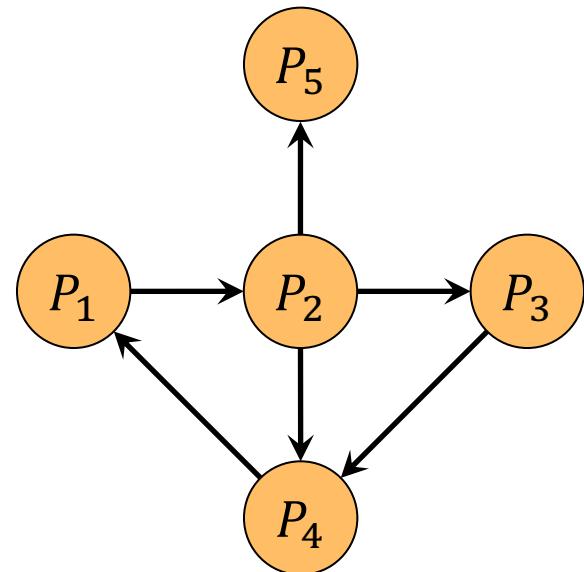
- If the graph contains no cycles, there is no deadlock.
  
- If the graph contains a cycle, then
  - If there is only one instance per resource type, there is a deadlock.
  - If there are several instances per resource type, there is a possibility of deadlock.

Deadlock Detection

# Single Instance of Each Resource Type



Resource-Allocation Graph



Corresponding wait-for graph

# Single Instance of Each Resource Type

- Maintain a wait-for graph
  - Nodes are processes.
  - $P_i \rightarrow P_k$  if  $P_i$  is waiting for a resource that is held by  $P_k$ .
- Periodically invoke an algorithm to search for a cycle in the graph.
  - If there is a cycle, there exists a deadlock.
- An algorithm to detect a cycle in a graph with  $n$  vertices requires an order of  $n^2$  operations.

# Solution for multiple instances of a resource type

- $\text{Available}_{[m]}$ 
  - $\text{Available}_k = x$  means that  $x$  instances of resource type  $R_k$  are available.
- $\text{Allocation}_{[n][m]}$ 
  - $\text{Allocation}_{i,k} = x$  means that process  $P_i$  has  $x$  instances of resource type  $R_k$  currently allocated to it.
- $\text{Request}_{[n][m]}$ 
  - $\text{Request}_{i,k} = x$  means that process  $P_i$  is requesting  $x$  more instances of resource type  $R_k$ .

# Detection Algorithm

- Let  $Work[m]$  and  $Finish[n]$  be initialized as
  - $Work \leftarrow Available$
  - $Finish_i \leftarrow (Allocation_i = 0)$
- While there is  $i \mid (\neg Finish_i) \wedge (Request_i \leq Work)$  do
  - $Work \leftarrow Work + Allocation_i$
  - $Finish_i \leftarrow true$
- If  $\neg Finish_i$  for some  $i$ , the system is in deadlock and every  $P_i$  for which  $\neg Finish_i$  is deadlocked.
- This algorithm requires  $O(m \times n^2)$  operations.

# Example of Detection Algorithm

- Five processes  $P_0$  through  $P_4$
- Three resource types:  $A$  (7 instances),  $B$  (2) and  $C$  (6).
- Snapshot at time  $t_0$ :

	<i>Allocation</i>		
	<i>A</i>	<i>B</i>	<i>C</i>
$P_0$	0	1	0
$P_1$	2	0	0
$P_2$	3	0	3
$P_3$	2	1	1
$P_4$	0	0	2

	<i>Request</i>		
	<i>A</i>	<i>B</i>	<i>C</i>
	0	0	0
	2	0	2
	0	0	0
	1	0	0
	0	0	2

	<i>Available</i>		
	<i>A</i>	<i>B</i>	<i>C</i>
	0	0	0

- No deadlock, because sequence  $\langle P_0, P_2, P_1, P_3, P_4 \rangle$  results in  $Finish_i$  being true for all  $i$ .

# Example of Detection Algorithm

- Same problem, but  $P_2$  now requests one additional instance of resource type  $C$ .
- Snapshot at time  $t_0$ :

	<i>Allocation</i>		
	<i>A</i>	<i>B</i>	<i>C</i>
$P_0$	0	1	0
$P_1$	2	0	0
$P_2$	3	0	3
$P_3$	2	1	1
$P_4$	0	0	2

	<i>Request</i>		
	<i>A</i>	<i>B</i>	<i>C</i>
	0	0	0
	2	0	2
	0	0	1
	1	0	0
	0	0	2

	<i>Available</i>		
	<i>A</i>	<i>B</i>	<i>C</i>
	0	0	0

# Example of Detection Algorithm

- Same problem, but  $P_2$  now requests one additional instance of resource type  $C$ .
- Snapshot at time  $t_0$ :

	<i>Allocation</i>		
	<i>A</i>	<i>B</i>	<i>C</i>
$P_0$	0	1	0
$P_1$	2	0	0
$P_2$	3	0	3
$P_3$	2	1	1
$P_4$	0	0	2

	<i>Request</i>		
	<i>A</i>	<i>B</i>	<i>C</i>
	0	0	0
	2	0	2
	0	0	1
	1	0	0
	0	0	2

	<i>Finish</i>		
	<i>A</i>	<i>B</i>	<i>C</i>

	<i>Available</i>		
	<i>A</i>	<i>B</i>	<i>C</i>
	0	0	0
	<i>Work</i>		
	<i>A</i>	<i>B</i>	<i>C</i>
	0	0	0

# Example of Detection Algorithm

- Same problem, but  $P_2$  now requests one additional instance of resource type  $C$ .
- Snapshot at time  $t_0$ :

	<i>Allocation</i>							<i>Request</i>						<i>Available</i>			
	<i>A</i>	<i>B</i>	<i>C</i>		<i>A</i>	<i>B</i>	<i>C</i>		<i>A</i>	<i>B</i>	<i>C</i>		<i>T</i>		<i>A</i>	<i>B</i>	<i>C</i>
$P_0$	0	1	0		0	0	0								0	0	0
$P_1$	2	0	0		2	0	2								0	0	0
$P_2$	3	0	3		0	0	1								0	0	0
$P_3$	2	1	1		1	0	0								0	0	0
$P_4$	0	0	2		0	0	2								0	1	0

	<i>Work</i>		
	<i>A</i>	<i>B</i>	<i>C</i>
	0	1	0

- The system is now in a deadlock involving processes  $P_1, P_2, P_3$  and  $P_4$ .

# Detection-Algorithm Usage

- When, and how often, to invoke depends on:
  - How often a deadlock is likely to occur?
  - How many processes will need to be rolled back?
    - One for each disjoint cycle
- If the detection algorithm is invoked arbitrarily, there may be many cycles in the resource graph.
  - It may be difficult or even impossible to tell which of the many deadlocked processes “caused” the deadlock.

# Recovery from Deadlock: Process Termination

- There are two options
  - Abort all deadlocked processes
  - Abort one process at a time until the deadlock ends.
- In which order should we choose to abort?
  - Priority of the process.
  - How long the process has run, and how long it will run.
  - Resources that the process has used.
  - Resources that the process needs to complete.
  - How many processes will have to be terminated.
  - Is the process interactive or batch?

# Recovery from Deadlock: Resource Preemption

- Selecting a victim
  - Minimize cost.
- Rollback
  - Return to some safe state, restart process from there.
- Starvation
  - Same process may always be picked as victim, so it is a good idea to include number of rollbacks in cost factor.