

T24

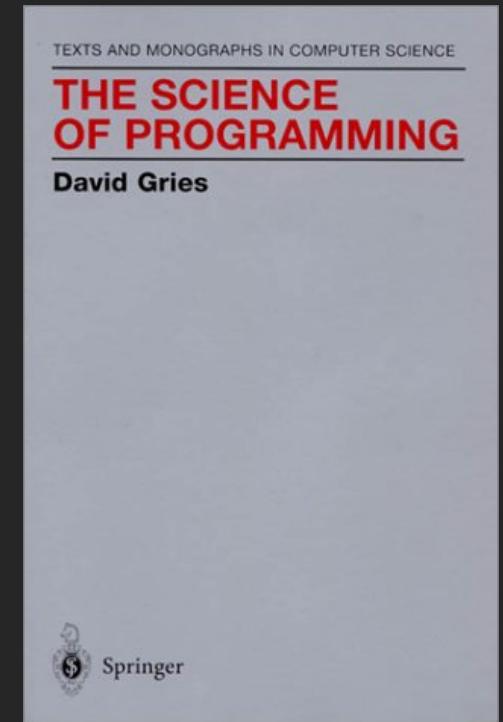
Event-Based Concurrency

Referência principal

Ch.33 of *Operating Systems: Three Easy Pieces* by Remzi and Andrea Arpaci-Dusseau (pages.cs.wisc.edu/~remzi/OSTEP/)

Discutido em classe em 24 de outubro de 2018

Modelos alternativos já eram assunto há 40 anos...



Dijkstra, E. W. 1975. Guarded commands, nondeterminacy and formal derivation of programs. *Commun. ACM* 18, 8 (Aug. 1975), 453-457.

Modelos alternativos já eram assunto há 40 anos...

On the Duality of Operating System Structures

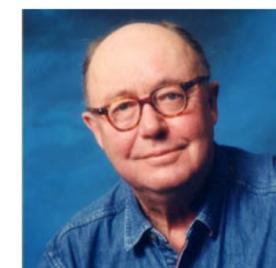
□ Hugh C. Lauer

- Adjunct Prof., Worcester Polytechnic Institute
- Xerox, Apollo Computer, Mitsubishi Electronic Research Lab, etc.
- Founded a number of businesses:
Real-Time Visualization unit of
Mitsubishi Electric Research Labs (MERL)



□ Roger M. Needham

- Prof., Cambridge University
- Microsoft Research, Cambridge Lab
- Kerberos, Needham-Schroeder security protocol, and key exchange systems



Modelos alternativos já eram assunto há 40 anos...

Message vs Procedure oriented systems (i.e. Events vs Threads)

- Are they really the same thing?
- Lauer and Needham show
 - 1) two models are duals
 - Mapping exists from one model to other
 - 2) dual programs are logically identical
 - Textually similar
 - 3) dual programs have identical performance
 - Measured in exec time, compute overhead, and queue/wait times

Modelos alternativos já eram assunto há 40 anos...

SEDA: An Architecture for Well-Conditioned, Scalable Internet Services (Welsh, 2001)

- 20 to 30 years later, still controversy!
- Analyzes threads vs event-based systems, finds problems with both
- Suggests trade-off: stage-driven architecture
- Evaluated for two applications
 - Easy to program and performs well

Modelos alternativos já eram assunto há 40 anos...

SEDA: An Architecture for Well-Conditioned, Scalable Internet Services (Welsh, 2001)

- Matt Welsh
 - Cornell undergraduate Alum (Worked on U-Net)
 - PhD from Berkeley (Worked on Ninja clustering)
 - Prof. at Harvard (Worked on sensor networks)
 - Currently at Google
- David Culler
 - Faculty at UC Berkeley
- Eric Brewer
 - Faculty at UC Berkeley (currently VP at Google)

Event-Based Concurrency

- As we have seen, it is not easy to generate a thread-safe concurrent application using the traditional procedural programming style.
 - For example, missing locks, deadlocks and other problems can easily arise.
- Event-based concurrency is another style of concurrent programming which has become popular in some modern systems.
 - For example, it is used in GUI-based applications, some types of internet servers, and server-side frameworks like **node.js**.
- **The idea is** to replace the complexities of multi-threading with **a deterministic, sequential model** that is easier to understand and debug, while giving the developer greater control over what is scheduled at a given moment in time.

Can we build concurrent systems without threads?

- We would like to allow a single-threaded program to cope with high-latency I/O devices by overlapping I/O with processing and other I/O.

Thus,

- Retaining control over concurrency
- Avoiding some of the problems common to multi-threaded applications
- How could we build a concurrent server like that?

The Basic Idea: An Event Loop

- The basic approach is called **event-based concurrency**:
 - Wait for something (i.e., an “**event**”) to occur.
 - When it does, check what type of event it is.
 - Do the small amount of work it requires.
- Example:

```
1. while (1) {  
2.     events = getEvents();  
3.     for (e in events)  
4.         processEvent(e);  
5. }
```



event handler

- While the event handler processes an event, this is the only activity taking place in the system.
- How can an event handler tell whether a message has arrived for it?

How to detect the occurrence of events

- Most systems provide a basic API, including `select()` or `poll()` system calls.
- They check whether there are incoming I/O requests to be serviced.
- They enable the developer to build a non-blocking event loop, which checks for incoming packets and reads from sockets with messages and replies as needed.
- With a single CPU and an event-based application, the problems found in concurrent programs are no longer present.
 - Since the event-based server cannot be interrupted by another thread, there is no need to acquire or release locks.
 - Thus, concurrency bugs common in multithreaded programs do not appear in the basic event-based approach.

Simple code using `select()`

```
1. int main(void) {
2.     // open and set up a bunch of sockets (not shown)
3.
4.     // main loop
5.     while (1) {
6.         // initialize the fd_set to all zero
7.         fd_set readFDs;
8.         FD_ZERO(&readFDs);
9.
10.        // now set the bits for the descriptors
11.        // this server is interested in
12.        // (for simplicity, all of them from min to max)
13.        for (int fd = minFD; fd < maxFD; fd++)
14.            FD_SET(fd, &readFDs);
15.
16.        // do the select
17.        int rc = select(maxFD + 1, &readFDs, NULL, NULL, NULL);
18.
19.        // check which actually have data using FD_ISSET()
20.        for (int fd = minFD; fd < maxFD; fd++)
21.            if (FD_ISSET(fd, &readFDs))
22.                processFD(fd);
23.    }
24. }
```

What if an event leads to a system call that might block?

- In a thread-based server, while the thread issuing the blocking I/O request suspends (waiting for the I/O to complete), other threads can run.
- This natural overlap of I/O and other operations is what makes thread-based programming quite natural and straightforward.
- With an event-based approach, however, there are no other threads to run: just the main event loop.
 - This implies that if an event handler issues a call that blocks, the entire server will also block until the call completes.
 - When the event loop blocks, the system sits idle, and thus is a huge potential waste of resources.
- We thus have a rule that must be obeyed in event-based systems: no blocking calls are allowed.

Why Simpler? No Locks Needed

- The event-based server cannot be interrupted by another thread.
 - With a single CPU and an event-based application.
 - It is decidedly **single threaded**.
 - Thus, *concurrency bugs* common in threaded programs **do not manifest** in the basic event-based approach.

A Problem: Blocking System Calls

- What if an event requires that you issue **a system call** that might block?
 - There are no other threads to run: *just the main event loop*
 - The entire server will do just that: **block until the call completes.**
 - Huge potential waste of resources

In event-based systems no blocking calls are allowed.

A Solution: Asynchronous I/O

- Enable an application to issue an I/O request and **return control immediately** to the caller, before the I/O has completed.
 - Example:

```
struct aiocb {  
    int aio_fildes; /* File descriptor */  
    off_t aio_offset; /* File offset */  
    volatile void *aio_buf; /* Location of buffer */  
    size_t aio_nbytes; /* Length of transfer */  
};
```

- An Interface provided on *macOS*
- The APIs revolve around a basic structure, the `struct aiocb` or **AIO control block** in common terminology.

A Solution: Asynchronous I/O (Cont.)

- Asynchronous API:
 - To issue an asynchronous read to a file

```
int aio_read(struct aiocb *aiocbp);
```
 - If successful, it returns right away and the application can continue with its work.
 - Checks whether the request referred to by aiocbp has completed.

```
int aio_error(const struct aiocb *aiocbp);
```
 - An application can **periodically poll** the system via aio_error().
 - If it has completed, returns success.
 - If not, EINPROGRESS is returned.

A Solution: Asynchronous I/O (Cont.)

- Interrupt
 - Remedy [the overhead to check](#) whether an I/O has completed
 - Using **UNIX signals** to inform applications when an asynchronous I/O completes.
 - Removing the need to *repeatedly ask the system.*

Another Problem: State Management

- The code of event-based approach is generally **more complicated** to write than *traditional thread-based* code.
 - It must package up some program state for the next event handler to use when the I/O completes.
 - The state the program needs is on the stack of the thread. → **manual stack management**

Another Problem: State Management (Cont.)

- **Example** (an event-based system):

```
int rc = read(fd, buffer, size);  
rc = write(sd, buffer, size);
```

- First **issue** the read asynchronously.
- Then, **periodically check** for completion of the read.
- That call informs us that the **read is complete**.
- How does the event-based server know **what to do?**

Another Problem: State Management (Cont.)

- **Solution:** continuation
 - **Record** the needed information to finish processing this event *in some data structure.*
 - When the event happens (i.e., when the disk I/O completes), **look up** the needed information and process the event.

What is still difficult with Events.

- Systems moved from a single CPU to multiple CPUs.
 - Some of the simplicity of the event-based approach disappeared.
- It does not integrate well with certain kinds of systems activity.
 - **Ex. Paging:** A server will not make progress until page fault completes (implicit blocking).
- Hard to manage overtime: The exact semantics of various routines changes.
- Asynchronous disk I/O never quite integrates with asynchronous network I/O in as simple and uniform a manner as you might think.

ASIDE: Unix Signals

- Provide a way to communicate with a process.
 - *HUP* (hang up), *INT*(interrupt), *SEGV*(segmentation violation), and etc.
 - **Example:** When your program encounters a *segmentation violation*, the OS sends it a *SIGSEGV*.

```
#include <stdio.h>
#include <signal.h>
void handle(int arg) {
    printf("stop wakin' me up...\n");
}

int main(int argc, char *argv[]) {
    signal(SIGHUP, handle);
    while (1)
        ; // doin' nothin' except catchin' some sigs
    return 0;
}
```

A simple program that goes into an infinite loop

ASIDE: Unix Signals (Cont.)

- You can send signals to it with the **kill command** line tool.
 - Doing so will *interrupt the main while loop* in the program and run the handler code `handle()`.

```
prompt> ./main &
[3] 36705
prompt> kill -HUP 36705
stop wakin' me up...
prompt> kill -HUP 36705
stop wakin' me up...
prompt> kill -HUP 36705
stop wakin' me up...
```