

# T31

# Locality and The Fast File System


*Referência principal*

Ch.41 of *Operating Systems: Three Easy Pieces* by Remzi and Andrea Arpaci-Dusseau ([pages.cs.wisc.edu/~remzi/OSTEP/](http://pages.cs.wisc.edu/~remzi/OSTEP/))

*Discutido em classe em 21 de novembro de 2018*

# Unix operating system

- Original Unix data structures on disk



- The good thing about the old file system
  - Simple and supported the basic abstractions, i.e. files and directories
  - Easy to use
  - A step-forward from earlier approaches
- The Problem
  - Terrible performance
    - Started off bad and got worse over time

# Issues of the old Unix file system

- Treated the disk as a random-access memory
  - Data was spread all over the disk without regard to positioning time
- File system easily fragmented due to uncaredful free space management
  - The policy was simply to take the next free block
    - For example, consider 4 files A, B, C and D, each with 2 blocks.



- If B and D are deleted, the layout becomes



- Now, if we create a 4-block file E we get



- And E is spread across the disk.

# Issues of the old Unix file system

- The original block size was too small (512 bytes)
  - Good because it reduced internal fragmentation
  - Bad because disk transfer was inefficient.

# How to organize on-disk data to improve performance?

What types of allocation policies are required?

How do we make the file system “disk aware”?

# Fast File System

- FFS was designed by a group at Berkeley in the early 80's.
- FFS structures and allocation policies were designed to be “disk aware” and improve performance.
  - It kept same API (`open()`, `read()`, `write()`, etc)
  - The internal implementation was heavily changed.

# Organizing Structure: The Cylinder Group


- FFS divides the disk into a bunch of groups. **(Cylinder Group)**
  - Modern file system call cylinder group as block group.

<b>G0</b>	<b>G1</b>	<b>G2</b>	<b>G3</b>	<b>G4</b>	<b>G5</b>	<b>G6</b>	<b>G7</b>	<b>G8</b>	<b>G9</b>
-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------	-----------

- These groups are used to improve seek performance.
  - By placing two files within the same group.
  - Accessing one after the other **will not be long seeks** across the disk.
  - FFS needs to allocate files and directories within each of these groups.

# Cylinder Groups


- FFS divides the disk into groups of  $N$  consecutive cylinders.



- Modern drives do not show such details about their geometry.
  - Instead, they export a logical address space of blocks and current file systems organize the drive into groups of consecutive blocks.



# Block Groups



- By placing two files within the same group, FFS ensures that accessing one after the other will not involve long seeks across the disk.
- To be able to place and manage files and directories into a group, FFS includes in it all the structures of a file system.



- Data structures for each block group
  - A copy of the **super block(S)** for reliability reasons
  - **inode bitmap(ib)** and **data bitmap(db)** to track free inodes and data blocks
  - **inodes** and **data blocks** are like those in the very-simple file system (VSFS)

# How To Allocate Files and Directories?

- Policy is “**keep related stuff together**”
  - But... what does “related” mean?
- The placement of directories...
  - Find a block group with a low number of allocated directories and a high number of free inodes.
  - Put the directory data and inode in that group.
- The placement of files...
  - Allocate data blocks of a file in the same block group of its inode
  - Place all files in the same block group as their directory

G0	G1	G2	G3	G4	G5	G6	G7	G8	G9
9 0		0 1		2 3		4 5		6 7	

# The Large-File Exception

- If the general policy of file placement is applied, a large file might
  - Entirely fill the block group it is first placed within
  - Prevent subsequent “related” files from being placed within this group and, thus, hurt file-access locality

Group 0	Group 1	Group 2	Group 3	Group 4	Group 5	Group 6	Group 7	Group 8	Group 9
		0 1 2 3 4 5 6 7 8 9							

- For large files, chunks are spread across the disk
  - Hurts performance, but this can be addressed by choosing chunk size
  - Amortization: reducing overhead by doing more work

Group 0	Group 1	Group 2	Group 3	Group 4	Group 5	Group 6	Group 7	Group 8	Group 9
6 7		0 1		4 5		2 3			8 9

# Amortization: How Big Do Chunks Have To Be?

- Estimating chunk size to achieve 50% of peak disk performance


- half of time seeking and
- half of time transferring

- Disk bandwidth: 40 MB/s

- Positioning time: 10ms


- How much data can this disk transfer in 10ms?

- $40 \frac{MB}{s} \times 0,01s = 0,4 MB = 0,4 \times 1024 KB = 409,6 KB$



# The Large-File Exception in FFS

- For large files, FFS took a simpler approach based on the inode structure
  - The first 12 direct blocks were placed in the same group as the inode
  - Each subsequent indirect block, and all the blocks it pointed to, were placed in a different block group.
  - With 4KB-blocks and 32-bit addresses, every 1024 blocks (4MB) of the file were in separate groups



# A few other things about FFS

- Internal fragmentation due to large (at that time!) block size
  - Solution: 512-byte sub-blocks
    - E.g. to create a file with 1 KB, use two sub-blocks, not an entire 4-KB block
- Parameterization
- Track buffers
- Long file names
  - Enabling more expressive names in the file system
- Symbolic links